

# CSCI446/946 Big Data Analytics

## Week 7      Advanced Analytical Theory and Methods: Classification

School of Computing and Information Technology  
University of Wollongong Australia

# Advanced Analytical Theory and Methods: **Classification**

- Overview of Classification
- **Decision Tree**
- **Naïve Bayes**
- Diagnostics of Classifier
- Additional Classification Models


All the figures, tables and codes are from the book “[Data Science and Big Data Analytics: Discovering, Analyzing, Visualizing and Presenting Data](#)” unless indicated otherwise.

# Overview of Classification

- Classification is a **fundamental** learning method that appears in applications related to data mining
- The primary task performed by classifiers is to **assign** class labels to **new** observations
- Classification methods are **supervised**
  - Start with a training set of **labelled** observations
  - Predict the outcome for **new** observations


# Decision Tree

- A decision tree uses **a tree structure** to specify sequences of decisions and consequences
- Given input variable  $X = \{x_1, x_2, \dots, x_n\}$ , the goal is to **predict** an output variable  $Y$




# Decision Tree

- Each **node** tests a particular input variable
- Each **branch** represents the decision made
- Classifying a new observation is to **traverse** this decision tree.



# Decision Tree

- The **depth** of a node is the **minimum** number of steps required to reach the node from root
- **Leaf** nodes are at the end of the last branches on the tree, representing class labels



# Decision Tree

- Use cases
  - **Classify animals**: questions (like cold-blooded or warm-blooded, mammal or not mammal) are answered to arrive at a certain classification
  - Checklist of symptoms during a doctor's **evaluation of a patient**
  - Retailers use decision tree to **predict response rates** to marketing and promotions
  - Financial institutions use it for **loan application**

# Decision Tree

- **An example:** A bank markets its term deposit product. So the bank needs to **predict** which clients would subscribe to a term deposit
  - The bank collects a dataset of 2000 **previous** clients with **known** “subscribe or not”.
  - **Input variables** to describe each client are
    - Job, marital status, education level, credit default, housing loan, personal loan, contact type, previous campaign contact



# Decision Tree

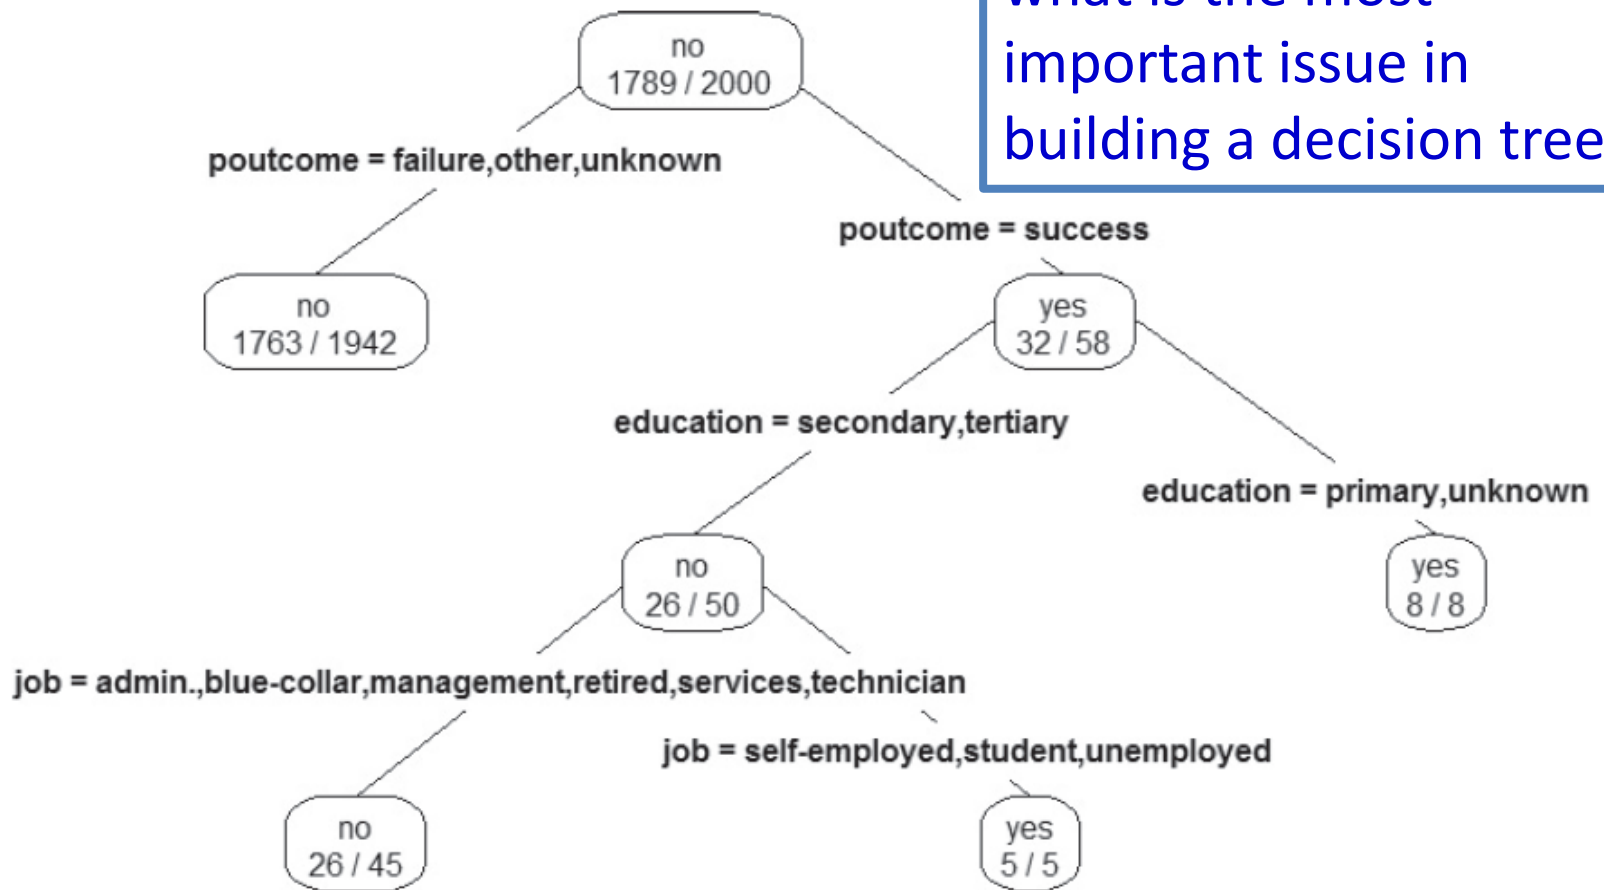
	job	marital	education	default	housing	loan	contact	poutcome	subscribed
1	management	single	tertiary	no	yes	no	cellular	unknown	no
2	entrepreneur	married	tertiary	no	yes	yes	cellular	unknown	no
3	services	divorced	secondary	no	no	no	cellular	unknown	yes
4	management	married	tertiary	no	yes	no	cellular	unknown	no
5	management	married	secondary	no	yes	no	unknown	unknown	no
6	management	single	tertiary	no	yes	no	unknown	unknown	no
7	entrepreneur	married	tertiary	no	yes	no	cellular	failure	yes
8	admin.	married	secondary	no	no	no	cellular	unknown	no
9	blue-collar	married	secondary	no	yes	no	cellular	other	no
10	management	married	tertiary	yes	no	no	cellular	unknown	no
11	blue-collar	married	secondary	no	yes	no	cellular	unknown	no
12	management	divorced	secondary	no	no	no	unknown	unknown	no
13	blue-collar	married	secondary	no	yes	no	cellular	unknown	no
14	retired	married	secondary	no	no	no	cellular	unknown	no
15	management	single	tertiary	no	yes	no	cellular	unknown	no

...

The training dataset of the bank example

# Decision Tree

From your point of view,  
what is the most  
important issue in  
building a decision tree?



A decision tree built over the bank marketing training dataset

# The General Algorithm of DT


- The **objective** of a decision tree algorithm
  - Construct a tree **T** from a training set **S**
- The algorithm picks the **most informative attribute** to branch the tree and does this **recursively** for each of the sub-trees.
- The most informative attribute is identified by
  - **Information gain**, calculated based on Entropy

# The General Algorithm of DT

- Entropy

Given a class  $X$  and its label  $x \in X$ , let  $P(x)$  be the probability of  $x$ .  $H_x$ , the entropy of  $X$ , is defined as

$$H_x = - \sum_{\forall x \in X} P(x) \log_2 P(x)$$



Question:

In the previous bank marketing dataset, there are 2000 customers in total. Among them, 1789 subscribed term deposit. What is the entropy of the output variable “subscribed” ( $H_{\text{subscribed}}$ )?

*Entropy of coin flips, where  $X=1$  represents heads*

# The General Algorithm of DT

- Conditional entropy

Given an attribute  $X$ , its value  $x$ , its outcome  $Y$ , and its value  $y$ , conditional entropy  $H_{Y|X}$  is the remaining entropy of  $Y$  given  $X$ ,

$$\begin{aligned} H_{Y|X} &= \sum_x P(x) H(Y|X=x) \\ &= - \sum_{\forall x \in X} P(x) \sum_{\forall y \in Y} P(y|x) \log_2 P(y|x) \end{aligned}$$

# The General Algorithm of DT

- Assume the attribute  $X$  is “contact”
  - Its value  $x$  takes one value in {cellular, telephone, unknown}
- The outcome  $Y$  is “subscribed”
  - Its value  $y$  takes one value in {no, yes}

	Cellular	Telephone	Unknown
$P(\text{contact})$	0.6435	0.0680	0.2885
$P(\text{subscribed=yes} \mid \text{contact})$	0.1399	0.0809	0.0347
$P(\text{subscribed=no} \mid \text{contact})$	0.8601	0.9192	0.9653

# The General Algorithm of DT

The conditional entropy of the *contact* attribute is computed as shown here.

$$\begin{aligned} H_{\text{subscribed}|\text{contact}} &= -[0.6435 \cdot (0.1399 \cdot \log_2 0.1399 + 0.8601 \cdot \log_2 0.8601) \\ &\quad + 0.0680 \cdot (0.0809 \cdot \log_2 0.0809 + 0.9192 \cdot \log_2 0.9192) \\ &\quad + 0.2885 \cdot (0.0347 \cdot \log_2 0.0347 + 0.9653 \cdot \log_2 0.9653)] \\ &= 0.4661 \end{aligned}$$
$$\begin{aligned} H_{Y|X} &= \sum_x P(x) H(Y|X=x) \\ &= - \sum_{\forall x \in X} P(x) \sum_{\forall y \in Y} P(y|x) \log_2 P(y|x) \end{aligned}$$

	Cellular	Telephone	Unknown
P(contact)	0.6435	0.0680	0.2885
P(subscribed=yes   contact)	0.1399	0.0809	0.0347
P(subscribed=no   contact)	0.8601	0.9192	0.9653

# The General Algorithm of DT

- Information gain

The information gain of an attribute  $A$  is defined as the difference between the base entropy and the conditional entropy of the attribute,

$$InfoGain_A = H_S - H_{S|A}$$

$$\begin{aligned} InfoGain_{contact} &= H_{subscribed} - H_{subscribed|contact} \\ &= 0.4862 - 0.4661 = 0.0201 \end{aligned}$$

- It compares
  - The degree of **purity** of the **parent** node **before** a split
  - The degree of **purity** of the **child** node **after** a split



# The General Algorithm of DT

- The algorithm **splits on the attribute** with the **largest information gain** at each round

Attribute	Information Gain
<i>poutcome</i>	0.0289
<i>contact</i>	0.0201
<i>housing</i>	0.0133
<i>job</i>	0.0101
<i>education</i>	0.0034
<i>marital</i>	0.0018
<i>loan</i>	0.0010
<i>default</i>	0.0005

# The General Algorithm of DT

- The algorithm constructs sub-trees recursively **until** one of the following criteria is met
  - All the leaf nodes in the tree satisfy the **minimum purity threshold** (i.e., are pure enough)
  - There is **no sufficient information gain** by splitting on more attribute (i.e., not worth anymore)
  - Any other stopping criterion is satisfied (such as the **maximum depth** of the tree)

# Decision Tree Algorithms

- Popular decision tree algorithms
  - ID3, C4.5 and CART

```
1  ID3 (A, P, T)
2    if  $T \in \phi$ 
3      return  $\phi$ 
4    if all records in T have the same value for P
5      return a single node with that value
6    if  $A \in \phi$ 
7      return a single node with the most frequent value of P in T
8    Compute information gain for each attribute in A relative to T
9    Pick attribute D with the largest gain
10   Let  $\{d_1, d_2 \dots d_m\}$  be the values of attribute D
11   Partition T into  $\{T_1, T_2 \dots T_m\}$  according to the values of D
12   return a tree with root D and branches labeled  $d_1, d_2 \dots d_m$ 
       going respectively to trees  $ID3(A - \{D\}, P, T_1)$ ,
        $ID3(A - \{D\}, P, T_2)$ , . . .  $ID3(A - \{D\}, P, T_m)$ 
```

# Evaluating a Decision Tree


- Decision tree uses **greedy** algorithms
  - It always chooses the option that seems **the best** available **at that moment**
  - However, the option may not be the best **overall** and this could cause **overfitting**
  - An **ensemble technique** can address this issue by combining multiple decision trees that use random splitting

# Evaluating a Decision Tree

- Ways to evaluate a decision tree
  - Evaluate whether the splits of the tree **make sense** and whether the decision rules are **sound** (say, with domain experts)
  - Having too many layers and obtaining nodes with few members might be signs of **overfitting**
  - Use standard diagnostics tools for classifiers

# Evaluating a Decision Tree

- **Overfitting** in decision tree
  - The lack of training data
  - The biased training data
  - Too many layers or nodes
- Avoid overfitting
  - **Stop growing** the tree **early** before all training data are perfectly classified
  - Grow the full tree and then **post-prune** the tree



# Properties of Decision Tree

- Computationally **inexpensive**, easy to classify
- Classification rules can be **understood**
- Handle both **numerical** and **categorical** input
- Handle variables that have a **nonlinear** effect on the outcome, better than linear models
- **Not** a good choice if there are many **irrelevant** input variables
  - Feature selection will be needed

# Decision Tree in R

- `rpart` library is for modelling decision tree
- `rpart.plot` enables the plotting of a tree
- An example
  - Predict `whether to play golf`
  - Input variables: `weather outlook, temperature, humidity, and wind`

```
install.packages("rpart.plot")    # install package rpart.plot  
library("rpart")                 # load libraries  
library("rpart.plot")
```



# Decision Tree in R

```
play_decision <- read.table("DTdata.csv",header=TRUE,sep=",")
play_decision
```

	Play	Outlook	Temperature	Humidity	Wind
1	yes	rainy	cool	normal	FALSE
2	no	rainy	cool	normal	TRUE
3	yes	overcast	hot	high	FALSE
4	no	sunny	mild	high	FALSE
5	yes	rainy	cool	normal	FALSE
6	yes	sunny	cool	normal	FALSE
7	yes	rainy	cool	normal	FALSE
8	yes	sunny	hot	normal	FALSE
9	yes	overcast	mild	high	TRUE
10	no	sunny	mild	high	TRUE

```
fit <- rpart(Play ~ Outlook + Temperature + Humidity + Wind,
            method="class",
            data=play_decision,
            control=rpart.control(minsplit=1),
            parms=list(split='information'))
```

# Decision Tree in R

```
summary(fit)
```

```
Call:
```

```
rpart(formula = Play ~ Outlook + Temperature + Humidity + Wind,  
      data = play_decision, method = "class",  
      parms = list(split = "information"),  
      control = rpart.control(minsplit = 1))  
n= 10
```

	CP	nsplit	rel error	xerror	xstd
1	0.3333333	0	1	1.000000	0.4830459
2	0.0100000	3	0	1.666667	0.5270463

```
Variable importance
```

Wind	Outlook	Temperature
51	29	20

```
Node number 1: 10 observations,      complexity param=0.3333333  
predicted class=yes expected loss=0.3 P(node) =1  
class counts:      3      7  
probabilities: 0.300 0.700  
left son=2 (3 obs) right son=3 (7 obs)
```

```
Primary splits:
```

```
Temperature splits as RRL,      improve=1.3282860, (0 missing)  
Wind      < 0.5 to the right, improve=1.3282860, (0 missing)  
Outlook   splits as RLL,      improve=0.8161371, (0 missing)  
Humidity  splits as LR,       improve=0.6326870, (0 missing)
```

```
Surrogate splits:
```

```
Wind < 0.5 to the right, agree=0.8, adj=0.333, (0 split)
```

```
Node number 2: 3 observations,      complexity param=0.3333333  
predicted class=no  expected loss=0.3333333 P(node) =0.3  
class counts:      2      1  
probabilities: 0.667 0.333
```

```
left son=4 (2 obs) right son=5 (1 obs)
```

```
Primary splits:
```

```
Outlook splits as R-L,      improve=1.9095430, (0 missing)  
Wind    < 0.5 to the left, improve=0.5232481, (0 missing)
```

```
Node number 3: 7 observations,      complexity param=0.3333333  
predicted class=yes expected loss=0.1428571 P(node) =0.7
```

```
class counts:      1      6
```

```
probabilities: 0.143 0.857
```

```
left son=6 (1 obs) right son=7 (6 obs)
```

```
Primary splits:
```

```
Wind      < 0.5 to the right, improve=2.8708140, (0 missing)  
Outlook   splits as RLR,      improve=0.6214736, (0 missing)  
Temperature splits as LR-,    improve=0.3688021, (0 missing)  
Humidity  splits as RL,       improve=0.1674470, (0 missing)
```

```
Node number 4: 2 observations  
predicted class=no  expected loss=0 P(node) =0.2  
class counts:      2      0  
probabilities: 1.000 0.000
```


```
Node number 5: 1 observations  
predicted class=yes expected loss=0 P(node) =0.1  
class counts:      0      1  
probabilities: 0.000 1.000
```

```
Node number 6: 1 observations  
predicted class=no  expected loss=0 P(node) =0.1  
class counts:      1      0  
probabilities: 1.000 0.000
```

```
Node number 7: 6 observations  
predicted class=yes expected loss=0 P(node) =0.6  
class counts:      0      6  
probabilities: 0.000 1.000
```

# Decision Tree in R

```
rpart.plot(fit, type=4, extra=1)
```



# Decision Tree in R

- Prediction outcome for a new observation

```
newdata <- data.frame(Outlook="rainy", Temperature="mild",  
                      Humidity="high", Wind=FALSE)
```

```
predict(object, newdata = list(),  
        type = c("vector", "prob", "class", "matrix"))
```



```
predict(fit,newdata=newdata,type="prob")
```

```
no yes  
1  1  0
```



```
predict(fit,newdata=newdata,type="class")
```

```
1  
no  
Levels: no yes
```

# Summary on Decision Tree



Copyright © Alex Alexeev

From <http://www.projectdecisions.org/index-cartoon-decisiontree.html>

# Naïve Bayes Classifier

- A **probabilistic** classification method based on **Bayes'** theorem
- A naïve Bayes classifier **assumes** that the presence or absence of a particular feature of a class is **unrelated** to the presence or absence of other features (**conditional independence assumption**)
- Output includes a **class label** and its corresponding **probability score**

# Bayes' Theorem



$C$  is the class label  $C \in \{c_1, c_2, \dots, c_n\}$

$A$  is the observed attributes  $A = \{a_1, a_2, \dots, a_m\}$

$$P(C|A) = \frac{P(A|C) \cdot P(C)}{P(A)}$$

Posteriori probability =  $\frac{\text{likelihood} \cdot \text{priori probability}}{\text{evidence}}$

# Bayes' Theorem

- Two examples to understand this theorem

An example better illustrates the use of Bayes' theorem. John flies frequently and likes to upgrade his seat to first class. He has determined that if he checks in for his flight at least two hours early, the probability that he will get an upgrade is 0.75; otherwise, the probability that he will get an upgrade is 0.35. With his busy schedule, he checks in at least two hours before his flight only 40% of the time. Suppose John did not receive an upgrade on his most recent attempt. What is the probability that he did not arrive two hours early?

Another example involves computing the probability that a patient carries a disease based on the result of a lab test. Assume that a patient named Mary took a lab test for a certain disease and the result came back positive. The test returns a positive result in 95% of the cases in which the disease is actually present, and it returns a positive result in 6% of the cases in which the disease is not present. Furthermore, 1% of the entire population has this disease. What is the probability that Mary actually has the disease, given that the test is positive?



# Bayes' Theorem

- A more practical form of Bayes' theorem

$$P(c_i|A) = \frac{P(a_1, a_2, \dots, a_m | c_i) \cdot P(c_i)}{P(a_1, a_2, \dots, a_m)}, i = 1, 2, \dots, n$$

C is the class label  $C \in \{c_1, c_2, \dots, c_n\}$

A is the observed attributes  $A = \{a_1, a_2, \dots, a_m\}$

- Given  $A$ , how to calculate  $P(c_i|A)$ ?

# Naïve Bayes Classifier

- With two simplifications, Bayes' theorem induces a Naïve Bayes classifier
- First, **Conditional independence assumption**
  - Each attribute is **conditionally** independent of every other attribute **given** a class label  $c_i$

$$P(a_1, a_2, \dots, a_m | c_i) = P(a_1 | c_i) P(a_2 | c_i) \cdots P(a_m | c_i) = \prod_{j=1}^m P(a_j | c_i)$$

- This simplifies the computation of  $P(A | c_i)$

# Naïve Bayes Classifier

- Second, **ignore** the denominator  $P(A)$ 
  - Removing the denominator has no impact on the **relative** probability scores
- In this way, the classifier becomes

$$P(c_i|A) \propto P(c_i) \cdot \prod_{j=1}^m P(a_j|c_i) \quad i = 1, 2, \dots, n$$

$$P(c_i|A) \propto \log P(c_i) + \sum_{j=1}^m \log P(a_j|c_i) \quad i = 1, 2, \dots, n$$

# Naïve Bayes Classifier

- An example
  - With the bank marketing dataset, use Naïve Bayes Classifier to **predict** if a client would subscribe to a term deposit
- **Building** a Naïve Bayes classifier requires to calculate some **statistics** from **training** dataset
  - $P(A/c_i)$  for each class  $i=1,2,\dots,n$
  - $P(a_j/c_i)$  for each attribute  $j=1,2,\dots,m$  in each class

$$P(c_i|A) \propto P(c_i) \cdot \prod_{j=1}^m P(a_j|c_i) \quad i=1,2,\dots,n$$


# Naïve Bayes Classifier

- $P(A/c_i)$  for each class

$$P(\text{subscribed} = \text{yes}) \approx 0.11 \text{ and } P(\text{subscribed} = \text{no}) \approx 0.89$$

The training set contains several attributes: *job*,  
*marital*, *education*, *default*, *housing*, *loan*, *contact*, and *outcome*.

- $P(a_j/c_i)$  for each attribute  
in each class



$$P(\text{single} \mid \text{subscribed} = \text{yes}) \approx 0.35$$

$$P(\text{married} \mid \text{subscribed} = \text{yes}) \approx 0.53$$

$$P(\text{divorced} \mid \text{subscribed} = \text{yes}) \approx 0.12$$

$$P(\text{single} \mid \text{subscribed} = \text{no}) \approx 0.28$$

$$P(\text{married} \mid \text{subscribed} = \text{no}) \approx 0.61$$

$$P(\text{divorced} \mid \text{subscribed} = \text{no}) \approx 0.11$$

# Naïve Bayes Classifier

- Testing a Naïve Bayes classifier on a new data

$j$	$a_j$	$P(a_j \mid \text{subscribed} = \text{yes})$	$P(a_j \mid \text{subscribed} = \text{no})$
1	job = management	0.22	0.21
2	marital = married	0.53	0.61
3	education = secondary	0.46	0.51
4	default = no	0.99	0.98
5	housing = yes	0.35	0.57
6	loan = no	0.90	0.85
7	contact = cellular	0.85	0.62
8	outcome = success	0.15	0.01

$$P(\text{yes}|A) \propto 0.11 \cdot (0.22 \cdot 0.53 \cdot 0.46 \cdot 0.99 \cdot 0.35 \cdot 0.90 \cdot 0.85 \cdot 0.15) \approx 0.00023$$

$$P(\text{no}|A) \propto 0.89 \cdot (0.21 \cdot 0.61 \cdot 0.51 \cdot 0.98 \cdot 0.57 \cdot 0.85 \cdot 0.62 \cdot 0.01) \approx 0.00017$$

# Naïve Bayes Classifier

- An issue on **rare** event
  - What if one of the attribute values does NOT appear in a class  $c_i$  in a training dataset?
  - $P(a_j/c_i)$  for this attribute value will equal **zero**!
  - $P(c_i/A)$  will simply become **zero**!
- **Smoothing** technique
  - It assigns a small **nonzero** probability to **rare** events not included in a training dataset

# Naïve Bayes Classifier

- Laplace smoothing (add-one smoothing)
  - It pretends to see every outcome once more than it actually appears

$$P^*(x) = \frac{\text{count}(x) + 1}{\sum_x [\text{count}(x) + 1]}$$

$$P'(\text{single} \mid \text{subscribed} = \text{yes}) = (20 + 1) / [(20 + 1) + (70 + 1) + (10 + 1)]$$

$$P^{**}(x) = \frac{\text{count}(x) + \varepsilon}{\sum_x [\text{count}(x) + \varepsilon]} \quad \varepsilon \in [0, 1]$$



# Naïve Bayes Classifier

- Advantages
  - Simple to implement, commonly used for text classification
  - Handle high-dimensional data efficiently
  - Robust to overfitting with smoothing technique
- Disadvantages
  - Sensitive to correlated variables (Why?)
  - Not reliable for probability estimation

# Naïve Bayes in R

- Two methods
  - Build the classifier **from the scratch**
  - Call **naiveBayes** function from **e1071** package

```
install.packages("e1071")    # install package e1071
library(e1071)               # load the library
```

```
# read the data into a table from the file
sample <- read.table("sample1.csv", header=TRUE, sep=",")
# define the data frames for the NB classifier
traindata <- as.data.frame(sample[1:14,])
testdata <- as.data.frame(sample[15,])
```

# Naïve Bayes in R

```
model <- naiveBayes(Enrolls ~ Age+Income+JobSatisfaction+Desire,  
                    traindata)  
  
# display model  
model  
  
# predict with testdata  
results <- predict (model,testdata)  
# display results  
results  
[1] Yes  
Levels:  No Yes
```

```
# use the NB classifier with Laplace smoothing  
model1 = naiveBayes(Enrolls ~., traindata, laplace=.01)
```

# Diagnostics of Classifiers

- Confusion matrix

		Predicted Class	
		Positive	Negative
Actual Class	Positive	True Positives (TP)	False Negatives (FN)
	Negative	False Positives (FP)	True Negatives (TN)

		Predicted Class		Total
		Subscribe	Not Subscribed	
Actual Class	Subscribed	3	8	11
	Not Subscribed	2	87	89
Total		5	95	100

# Diagnostics of Classifiers

- Metrics used to **evaluate** classifiers

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \times 100\%$$

$$FPR = \frac{FP}{FP + TN}$$

$$TPR = \frac{TP}{TP + FN}$$

$$FNR = \frac{FN}{TP + FN}$$

$$Precision = \frac{TP}{TP + FP}$$

$$TPR \text{ (or Recall)} = \frac{TP}{TP + FN}$$

# Additional Classification Models

- Bagging
  - Bootstrap technique, **ensemble** method
- Boosting
  - Weighted combination, **ensemble** method
- Random Forest
  - Combination of decision trees, **ensemble** method
- Support Vector Machines
  - **Max-margin** linear classifier, **kernel** trick

# Summary

- Decision trees and Naïve Bayes classifier

Concerns	Recommended Method(s)
Output of the classification should include class probabilities in addition to the class labels.	Logistic regression, decision tree
Analysts want to gain an insight into how the variables affect the model.	Logistic regression, decision tree
The problem is high dimensional.	Naïve Bayes
Some of the input variables might be correlated.	Logistic regression, decision tree
Some of the input variables might be irrelevant.	Decision tree, naïve Bayes
The data contains categorical variables with a large number of levels.	Decision tree, naïve Bayes
The data contains mixed variable types.	Logistic regression, decision tree
There is nonlinear data or discontinuities in the input variables that would affect the output.	Decision tree

