

# **CSCI446/946 Big Data Analytics**

## **Week 11    Advanced Analytical Theory and Methods: Image Analysis**

School of Computing and Information Technology  
University of Wollongong Australia

# Advanced Analytical Theory and Methods: Image Analysis

- Overview of Image Analysis
- Collecting and Representing Image
- Image Recognition
- Bag-of-Visual-Words model
- Deep Convolutional Neural Networks



# Overview of Image Analysis

- Image analysis
  - Refers to the **representation**, **processing**, and **modelling** of visual data to derive useful insights
  - Suffers from the **semantic gap**
  - Visual data (image, video, ...) is **unstructured**
- **Semantic gap**
  - The gap between **high-level concepts** used by human and the **low-level features** used by computer

# Overview of Image Analysis

- Image processing
  - Scaling, translation, rotation, ...
  - Filtering, enhancement, histogram equalisation, ...

noisy lena



Gaussian filter



median filter

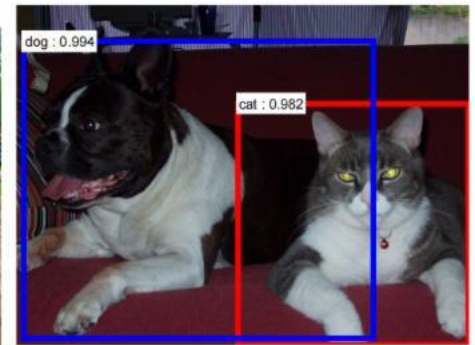



Wiener filter



# Overview of Image Analysis

- Image recognition (in a narrow sense)
  - Image classification
  - Object detection, localisation, tracking
  - Scene segmentation and reconstruction
  - Image search and retrieval



# Overview of Image Analysis

- Image classification



Face recognition



OCR recognition



Scene recognition

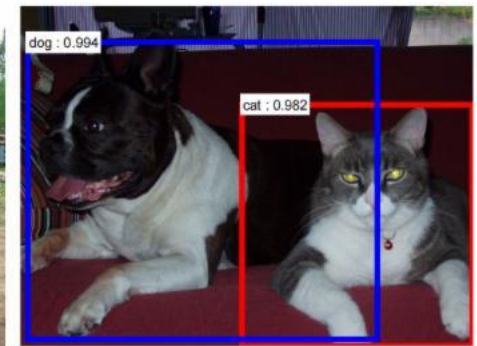



Object recognition

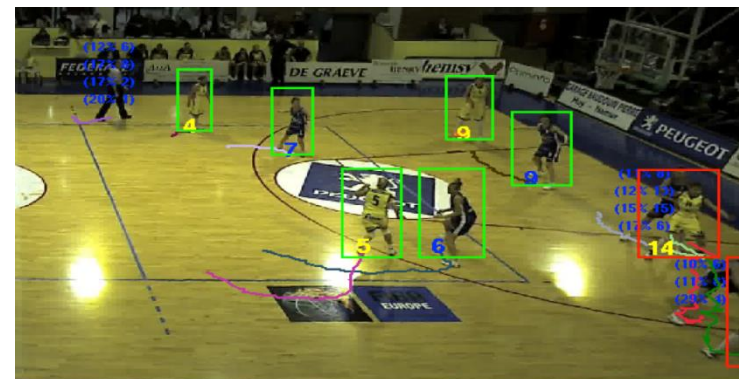



# Overview of Image Analysis

- Object detection, localisation, tracking



Object detection and localization





Object tracking (<https://www.youtube.com/watch?v=dKpRsdYSCLQ>)

# Overview of Image Analysis

- Scene segmentation and reconstruction



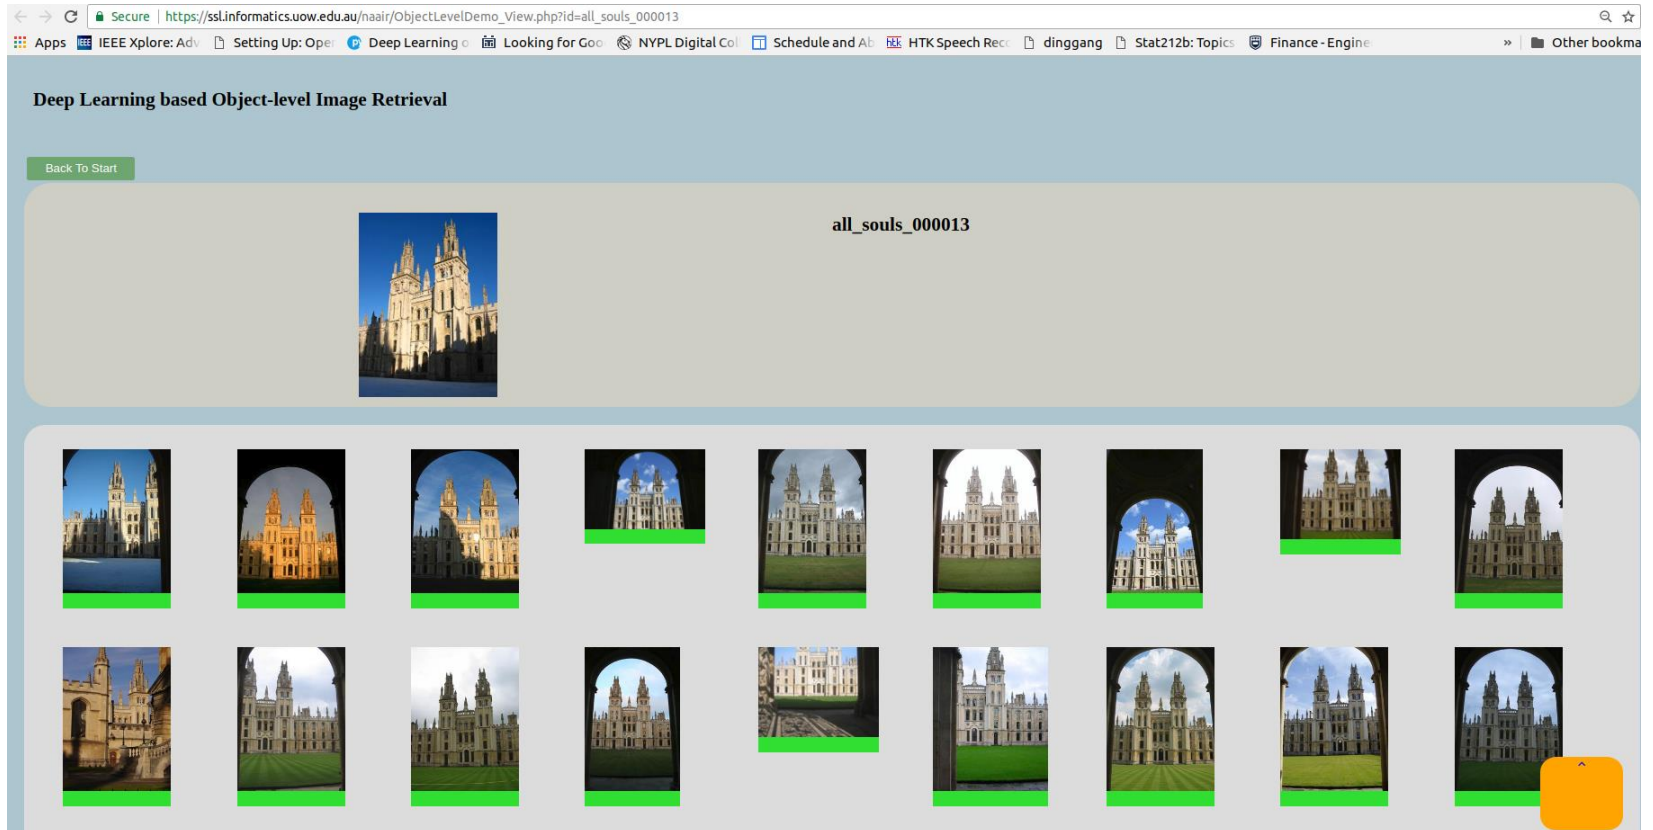
[Farabet et al. PAMI 2013]






# Overview of Image Analysis

- Image search and retrieval



# Overview of Image Analysis

- Image recognition benchmark datasets
  - [Caltech101 and Caltech256](#)
  - [PASCAL VOC project](#)
  - [ImageNet](#) The logo for ImageNet, featuring the word "IMAGENET" in a sans-serif font. The "I" is grey, "M" is grey, "A" is green, "G" is orange, "E" is red, "N" is grey, and "E" is grey.
  - [Microsoft COCO](#)
  - [ORL database of faces](#)
  - [Labelled Faces in the Wild](#)
  - A lot more...

# Image Analysis Steps

- Collection and labelling
  - Collect representative images from a given task and label the ground truth
- Image representation
  - Select and/or design appropriate image representations (**invariant** and **discriminative**)
- Image analysis techniques
  - Apply and/or design appropriate analysis techniques for the given tasks (classification, detection, tracking, segmentation, etc.)

# An Image Analysis Example

- A shop would like to **analyse customer shopping behaviour** via surveillance video
  - **Who** are their main customers?
  - **What** are their shopping patterns?



[https://en.wikipedia.org/wiki/Coles\\_Supermarkets](https://en.wikipedia.org/wiki/Coles_Supermarkets)



<http://www.dailytelegraph.com.au>

# Collecting Raw Image

- For image analysis, data must be collected before anything can happen
- The team
  - Starts by actively monitoring various sources for **representative images or videos**
  - Will deal with **unstructured data**
- Public **APIs** and Web scraper
- Be careful about the **rights** of the owner



# Representing Image


- Why representing images is **difficult**?
  - Scale, rotation, illumination, occlusion, background clutter, deformation, ...
  - **Invariant** and **Discriminative** representation

Cat:



# Representing Image

- **Traditional** representation (before year 2000)
  - Hand-crafted, global features
  - Intensity, colour, texture, shape, structure, etc.



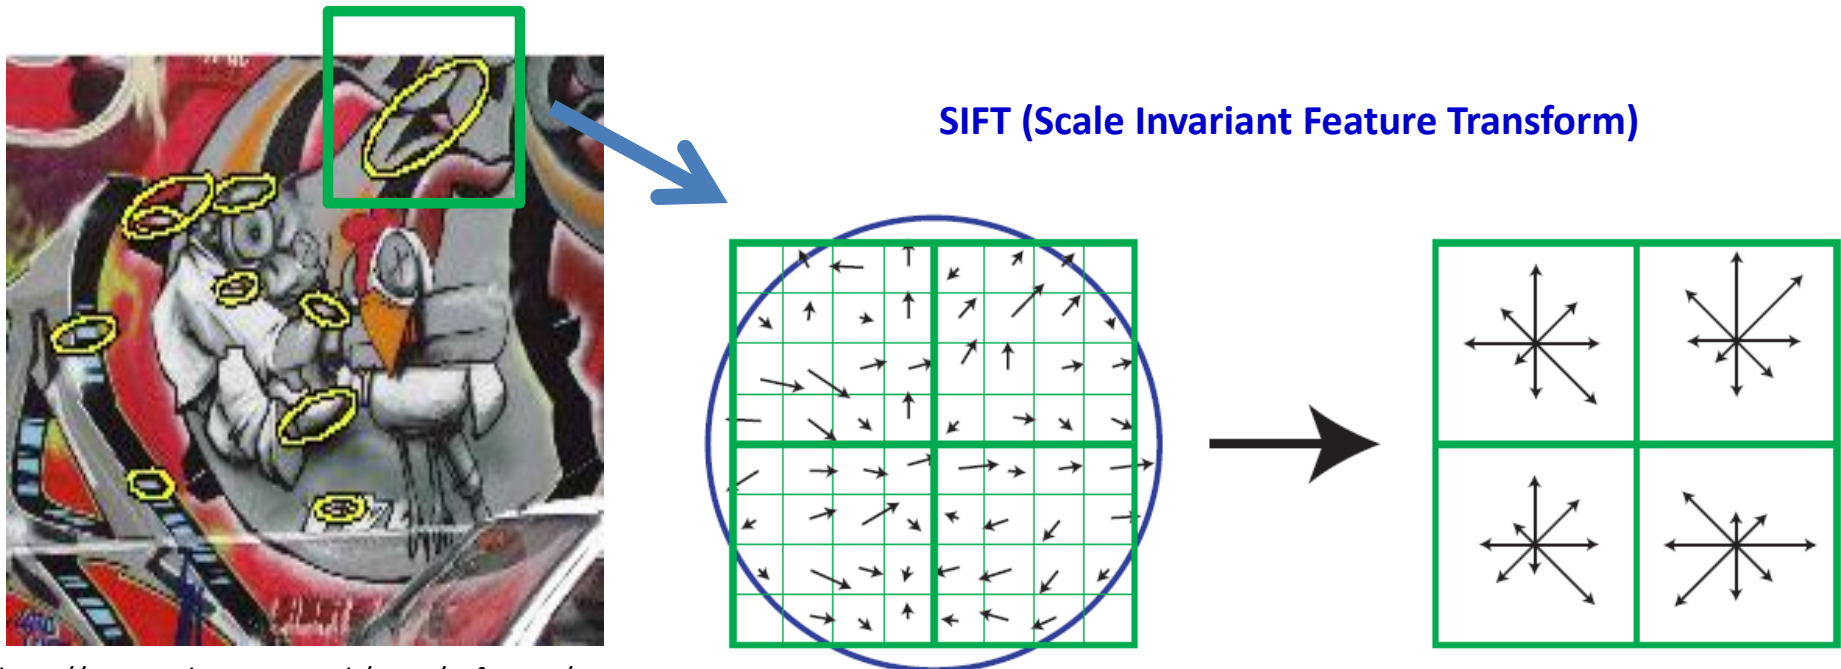
Colour histogram in a RGB space



Face recognition with raw pixel intensities

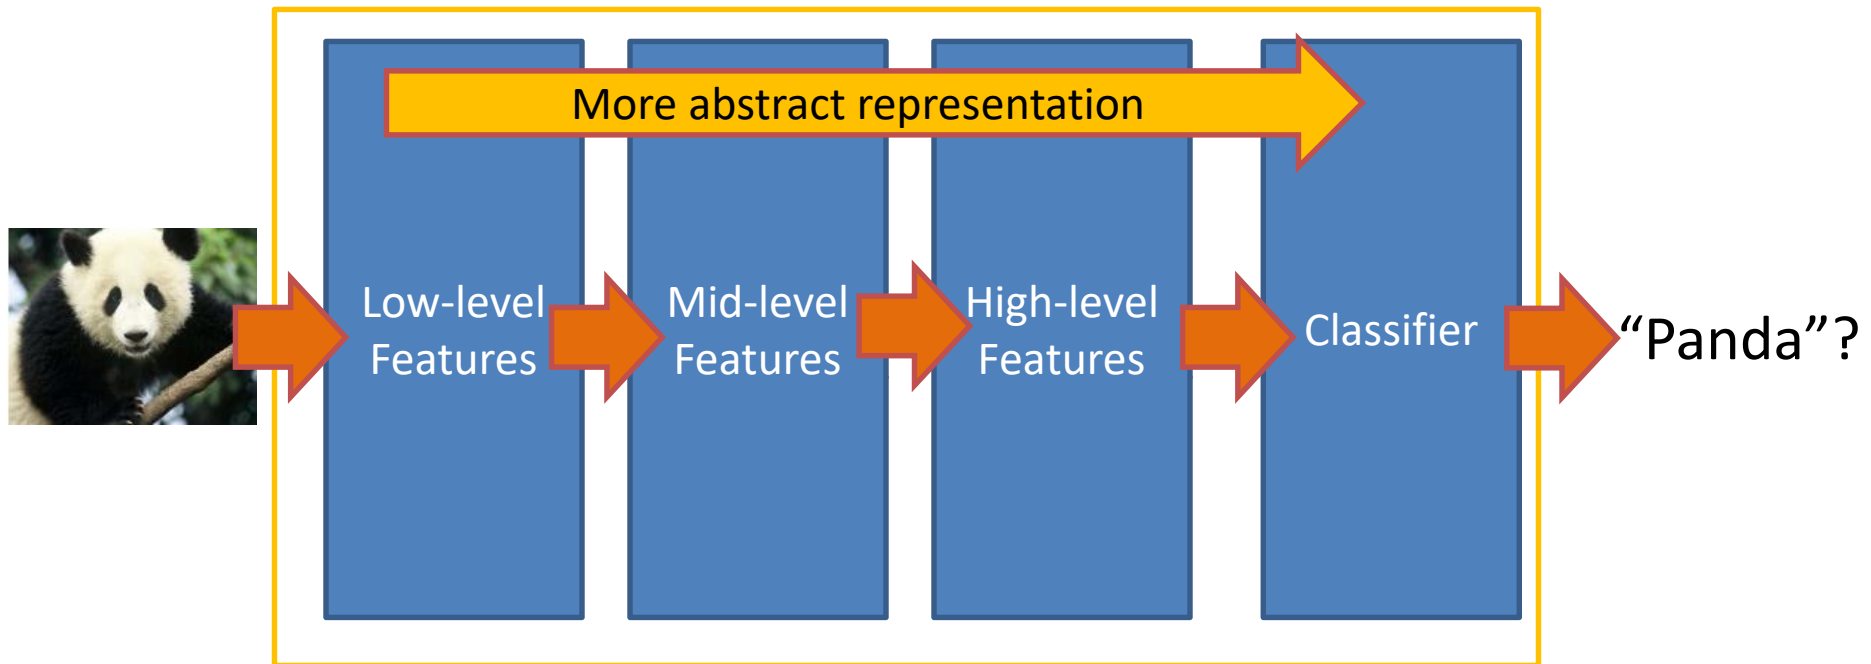
# Representing Image

- Days of the BoVW model (2000 ~ 2012)
  - **SIFT**, HOG, SURF, CENTRIST, filter-based, ...
  - **Invariant** to view angle, scale, illumination, ...



# Representing Image

- Era of Deep Learning (2012 ~ present)
  - Directly learn features representations from data




Deep Learning: train layers of features so that classifier works well.

Image courtesy of M. Ranzato

# Bag-of-Visual-Words Model



- Remember **the Bag-of-Words** representation?
  - A **document** becomes a **high-dimensional vector**, indicating the **presence/absence/frequency** of various **words** in this document





# Bag-of-Visual-Words Model



- BoVW model is borrowed from text analysis



# Bag-of-Visual-Words Model




Interest point detection  
or  
Dense sampling



The cropped detected regions


# Bag-of-Visual-Words Model

A close-up view






# Bag-of-Visual-Words Model



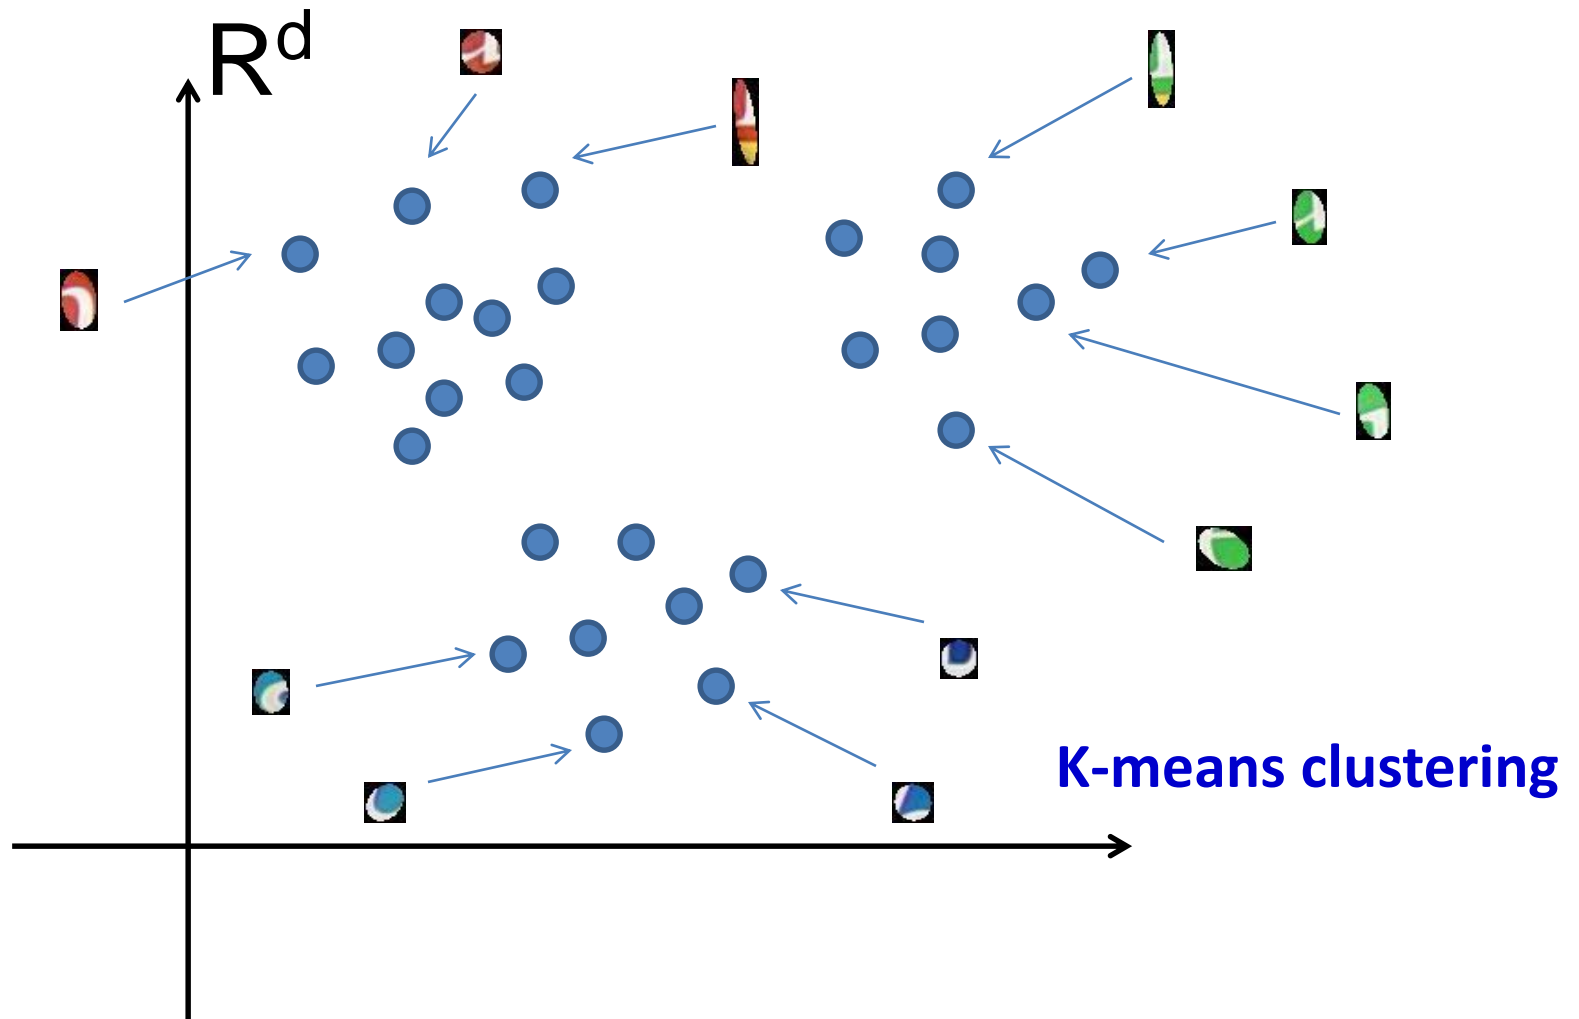
All images  
in a dataset



→ x 2  $R^d$

# Bag-of-Visual-Words Model


Cluster all features to generated “Visual Words”





# Bag-of-Visual-Words Model

Generated “Visual Words”




# Bag-of-Visual-Words Model

From an image to a **histogram**



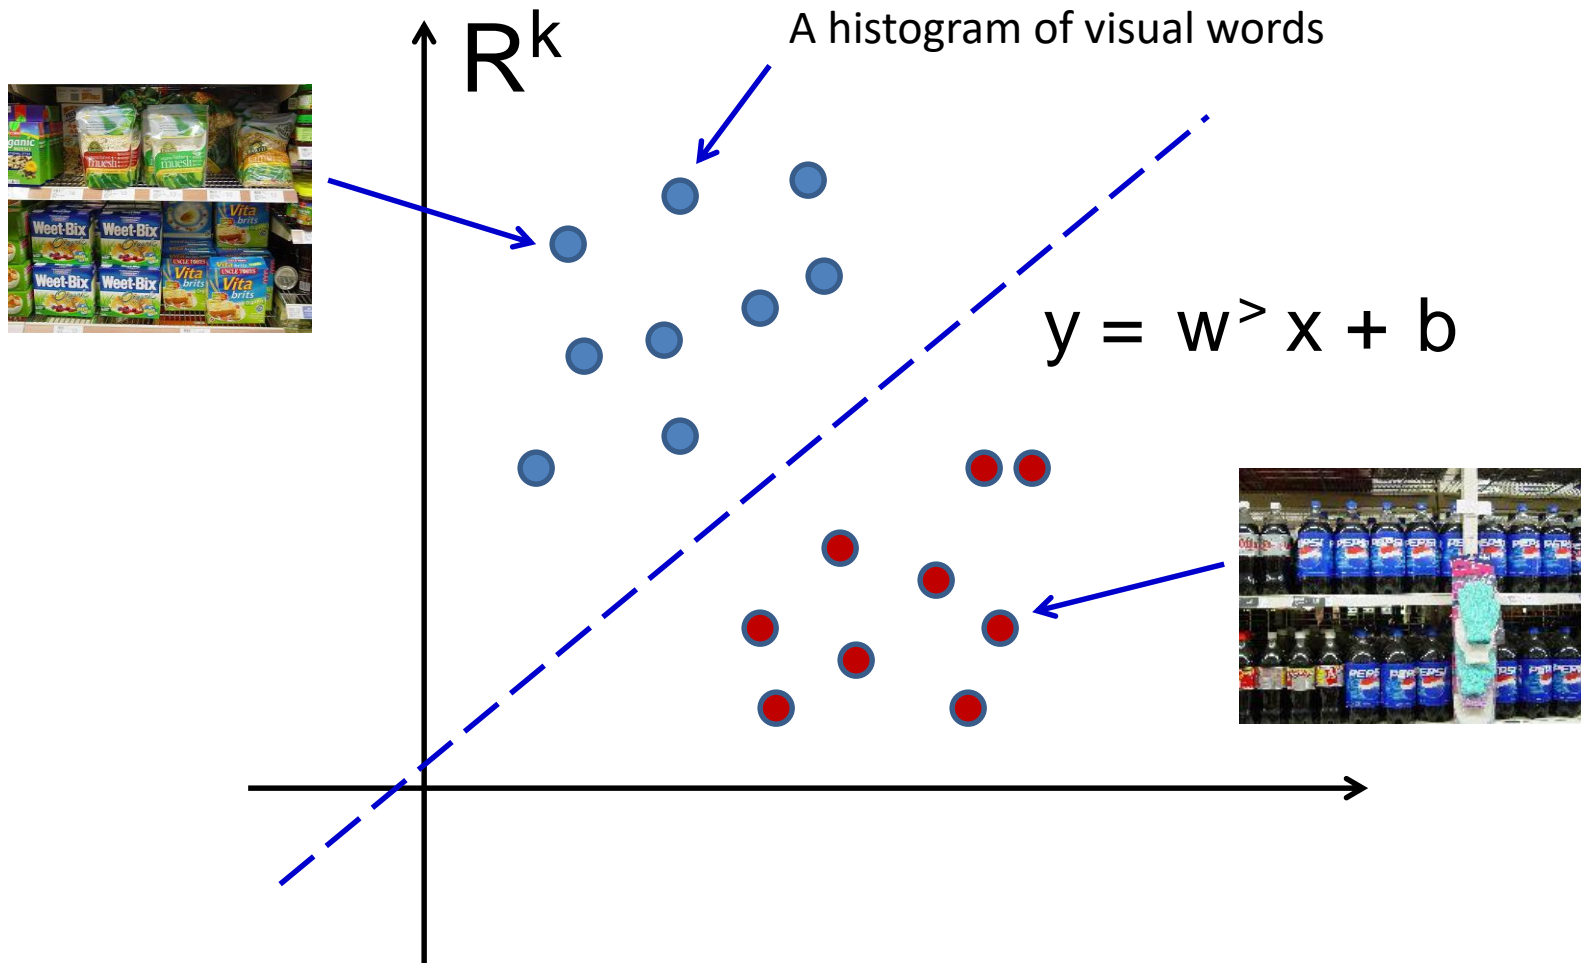
The number of occurrence of 1<sup>st</sup> “word” in this image

$[n_1, n_2, \dots, n_k] \in \mathbb{R}^k$




# Bag-of-Visual-Words Model

Classifying images



# Bag-of-Visual-Words Model

Procedure of the BoVW model based Image Recognition



# Deep Learning Model

MIT  
Technology  
Review

## 10 BREAKTHROUGH TECHNOLOGIES 2013

### Deep Learning

With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart. →

### Temporary Social Media

Messages that quickly self-destruct could enhance the privacy of online communications and make people freer to be spontaneous. →

### Prenatal DNA Sequencing

Reading the DNA of fetuses will be the next frontier of the genomic revolution. But do you really want to know about the genetic problems or musical aptitude of your unborn child? →

### Memory Implants

A maverick neuroscientist believes he has deciphered the code by which the brain forms long-term memories. Next: testing a prosthetic implant for people suffering from long-term memory loss. →

### Smart Watches

The designers of the Pebble watch realized that a mobile phone is more useful if you don't have to take it out of your pocket. →

### Ultra-Efficient Solar Power

Doubling the efficiency of a solar cell would completely change the economics of renewable energy. Nanotechnology just might make it possible. →

### Visual

- Images
- Videos

### Audio

- Speech
- Music

### Text

- Natural Language

### Planning

### Additive Manufacturing

Skeptical about 3-D printing? Get the world's largest manufacturer on the verge of using the technology to make jet parts. →

### Baxter: The Blue-Collar Robot

Rodney Brooks's newest creation is easy to interact with, but the complex innovations behind the robot show just how hard it is to get along with people. →

### Big Data from Cheap Phones

Collecting and analyzing information from cheap phones can provide surprising insights into how people move and behave – and... even help us understand the spread of diseases. →

### Supergrids

A new high-power circuit breaker could finally make highly efficient DC power grids practical. →

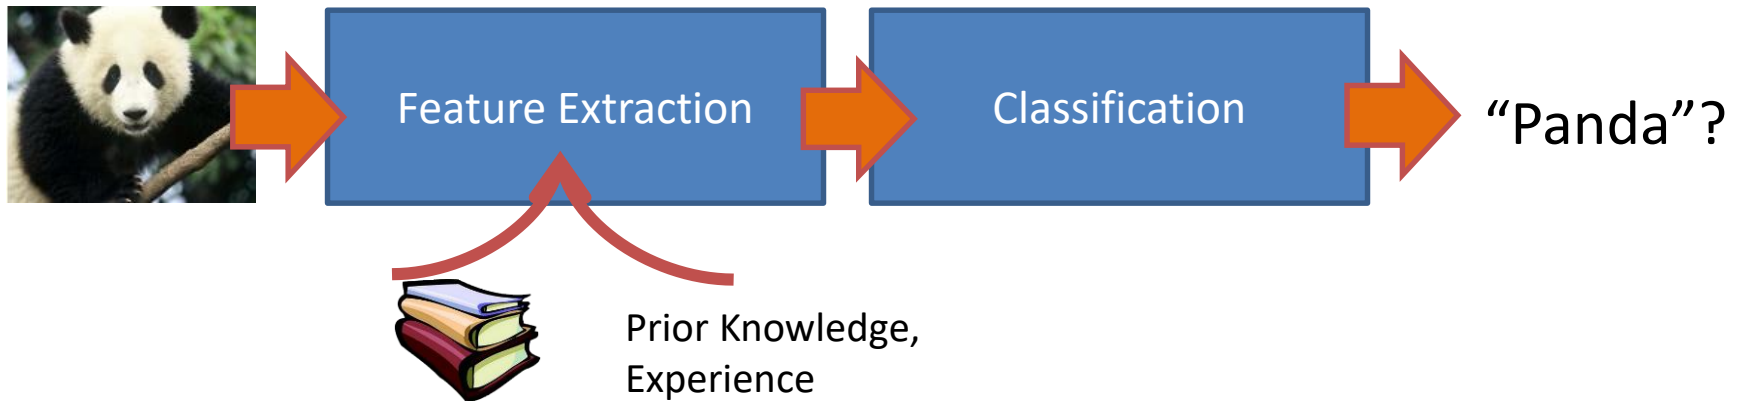


# Deep Learning Model

- Image Recognition
  - Faces, objects, poses, scenes, ...
- Video content analysis
  - Action, activities, events, summarization, ...
- Visual information management
  - Search, retrieval, indexing, browsing, ...
- Potential Outcome: AI
  - Computers can see and understand visual information
  - Robotics, self-driving cars, surveillance
  - ....

# Deep Learning Model

- We need **Invariant** and **discriminative** features!



Pose



Occlusion



Multiple objects

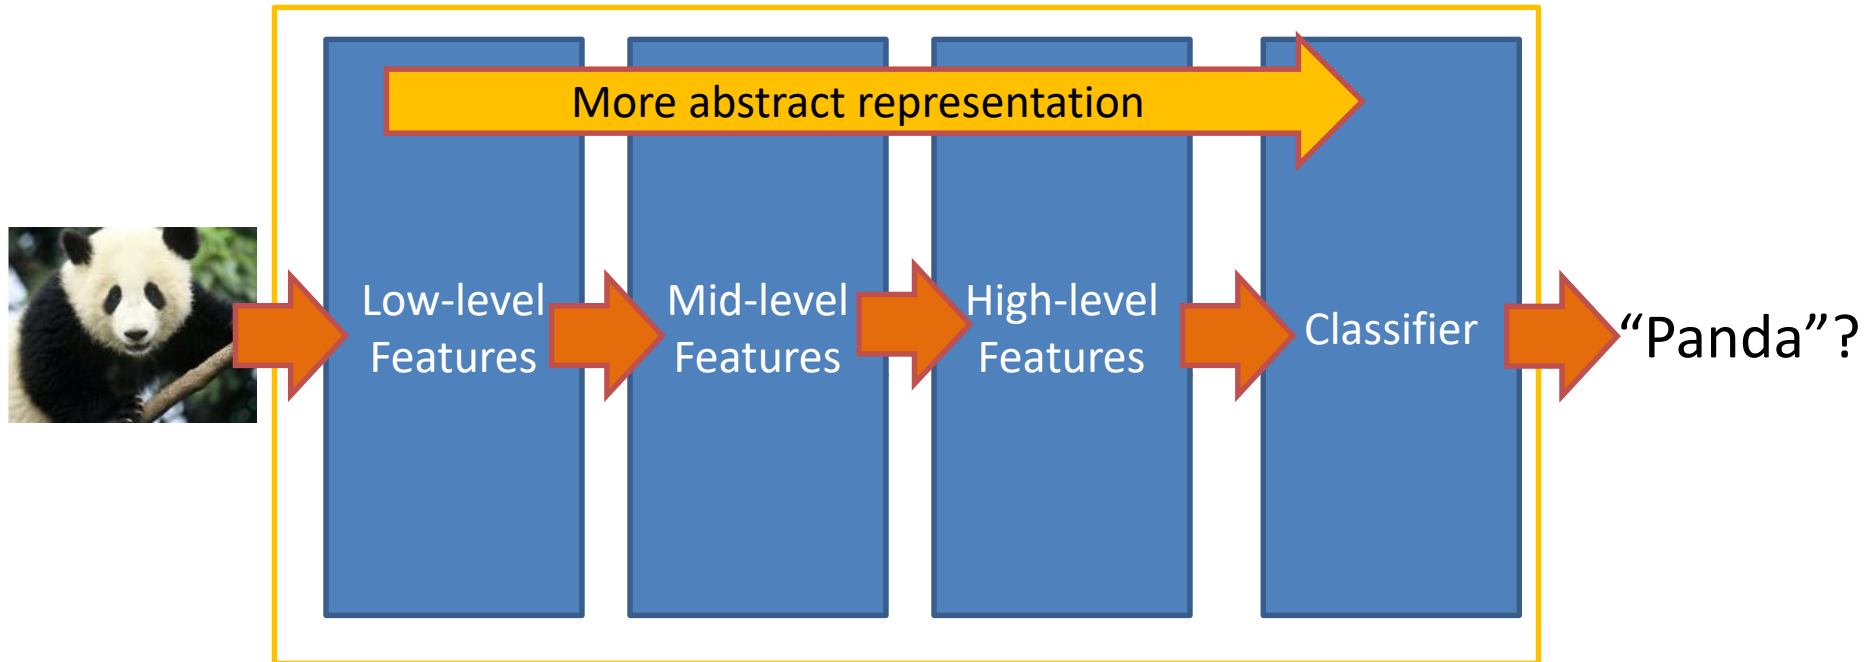


Inter-class similarity



# Deep Learning Model

- **Directly learn** features representations from data.
- **Joint learn** feature representation and classifier.





Deep Learning: train layers of features so that classifier works well.

Image courtesy of M. Ranzato

# Deep Learning Model

- Inspired by the way that human brain processes information
- Many **layers** of **non-linear** processing stages



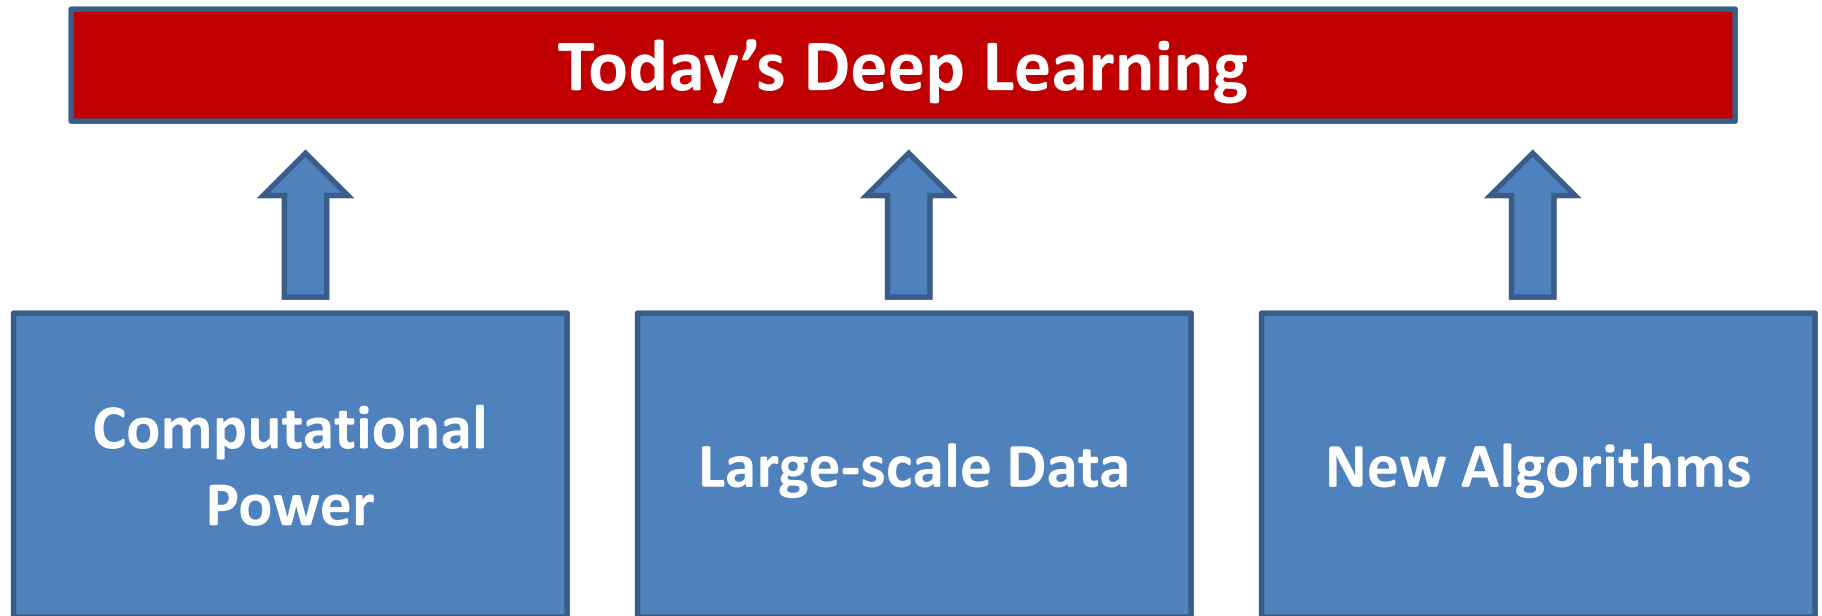
# Deep Learning Model

Have we been here before?

➤ Yes.

- Basic ideas common to past **neural networks** research
- Standard machine learning strategies still relevant.

➤ No.





# Deep Learning Model

## Convolutional Neural Networks (CNNs)

- A special **multi-stage architecture** inspired by visual system
  - Higher stages compute more global, more invariant features





Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

# Deep Learning Model

## CNNs: **ImageNet** Breakthrough



- Krizhevsky et al. win 2012 ImageNet classification with a **much bigger ConvNet**
  - **deeper**: 7 stages vs 3 before
  - **larger**: 60 million parameters vs 1 million before
  - **16.4%** error (top-5) vs Next best 26.2% error
- This was made possible by:
  - **fast hardware**: GPU-optimized code
  - **big dataset**: 1.2 million images vs thousands before
  - **better regularization**: dropout et al.

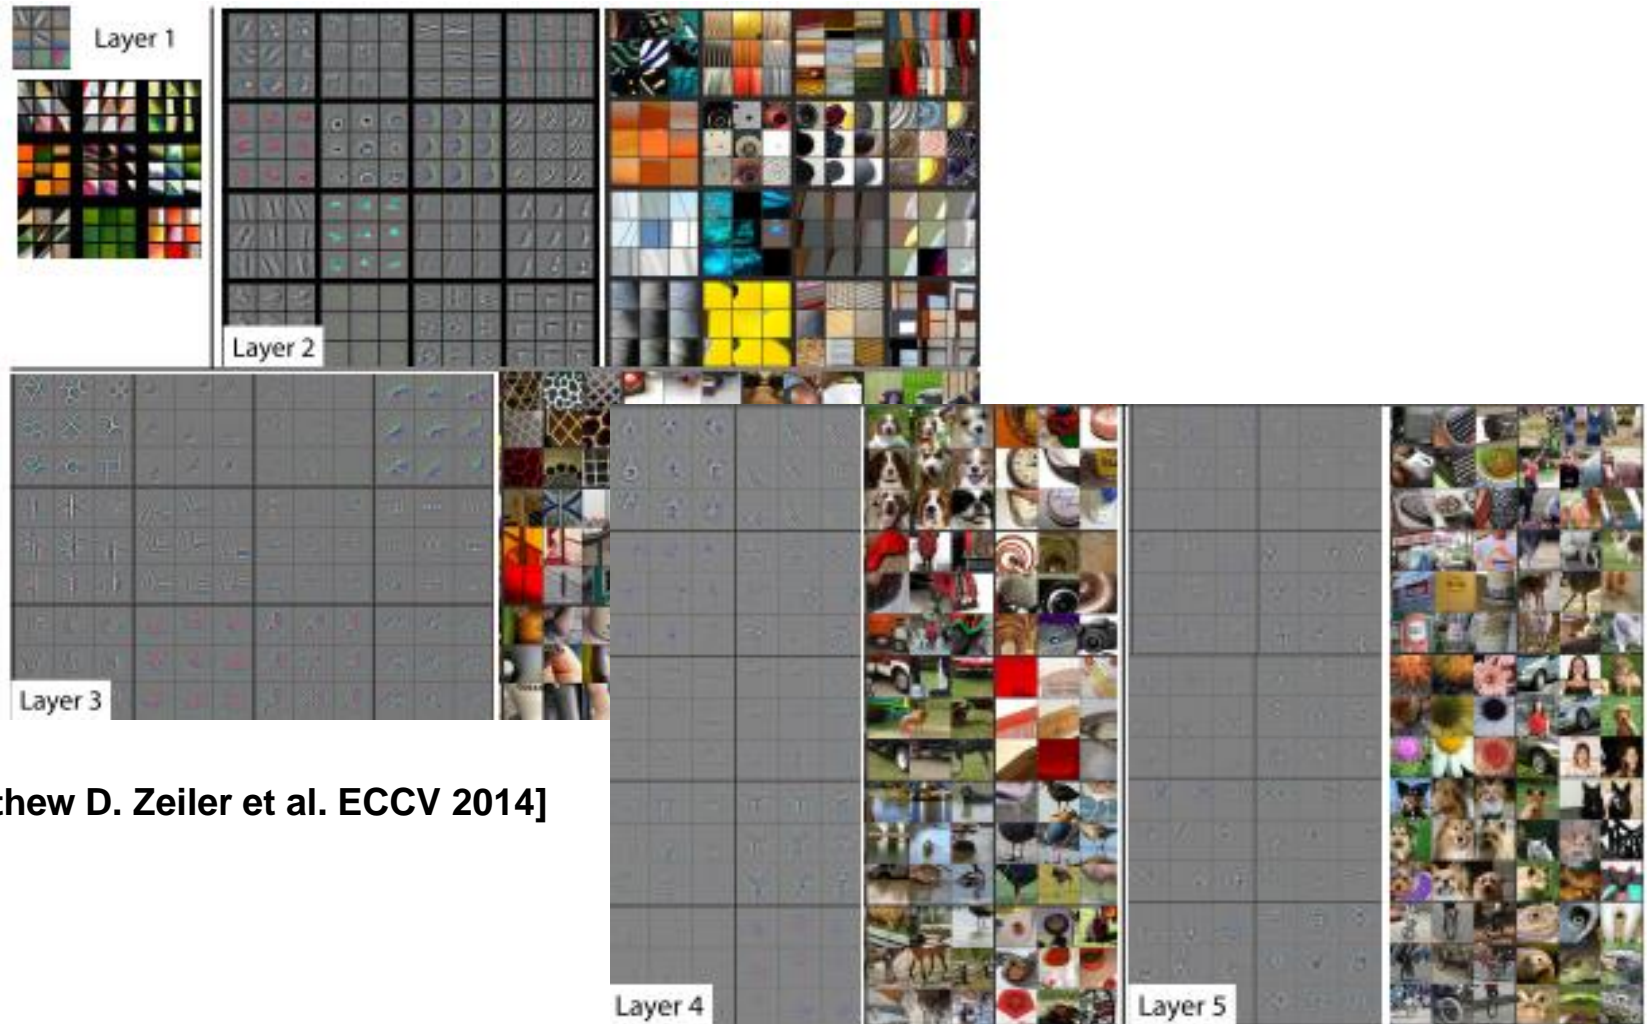
IMAGENET



Image courtesy of Deng et al.


# Deep Learning Model

## Learned Features of CNNs




[Matthew D. Zeiler et al. ECCV 2014]

# Deep Learning Model



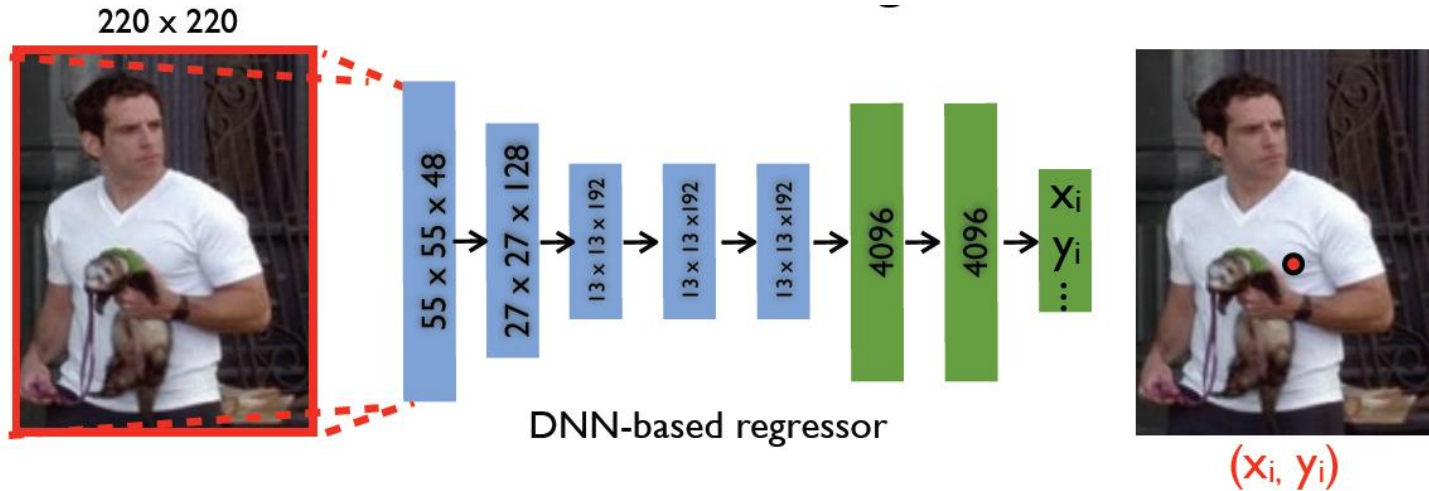
**Object detection** (Source: Rich feature hierarchies for accurate object detection and semantic segmentation, CVPR 2014)




**Face Recognition** (Source: DeepFace: Closing the Gap to Human-Level Performance in Face Verification, CVPR 2014)



# Deep Learning Model



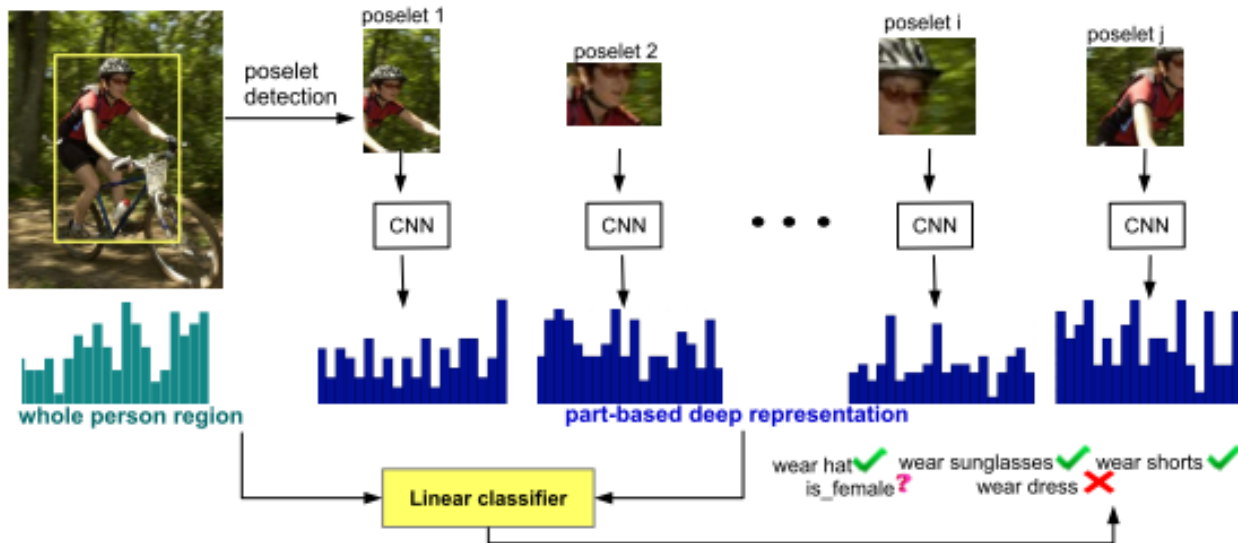
**Pose estimation** (DeepPose: Human Pose Estimation via Deep Neural Networks, CVPR2014)



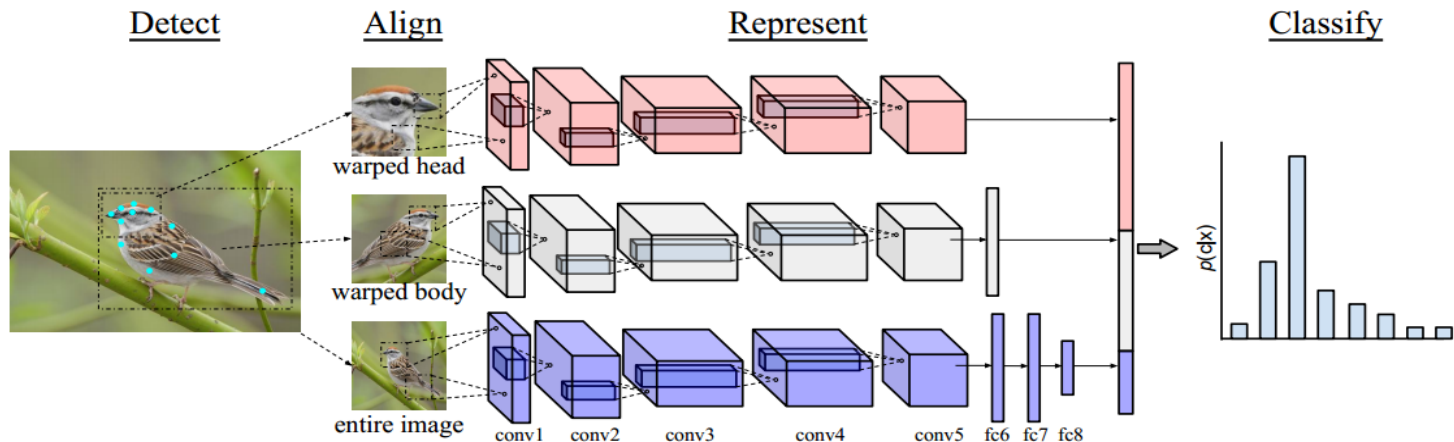
**Image Segmentation** (Source: SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation, IEEE TPAMI 2016)



# Deep Learning Model




Human attribute classification [Ning Zhang et al. CVPR 2014]

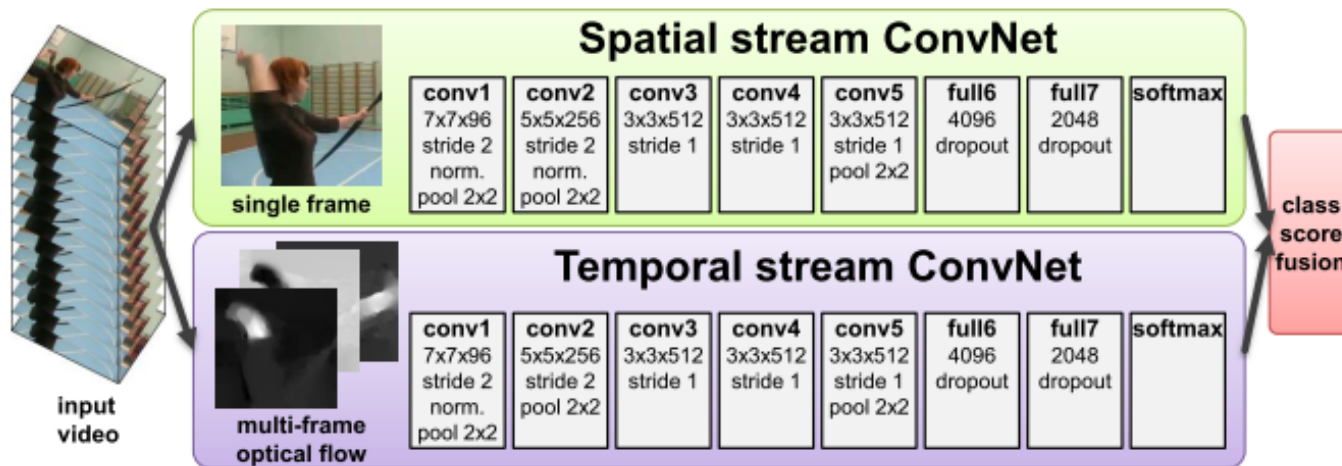


Fine-grained image recognition [Branson et al. arXiv 2014]

# Deep Learning Model



Large-scale Video Classification [Karpathy et al. CVPR 2014]



Action Recognition [Simonyan et al. arXiv 2014]

# Deep Learning Model

## Image Representation: From **SIFT** to **CNNs**

- Three main approaches
  - **Directly use pre-trained CNNs** models
    - to extract image feature representations
  - **Fine-tune pre-trained CNNs** models
    - with the images from recognition tasks
  - **BoVW model** based on CNN features
    - “**Deep SIFT**”

# Deep Learning Model

- Directly use pre-trained CNNs
  - **Which** layer to use?
  - How to **pool** the features in a **convolutional** layer?




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

# Deep Learning Model

- Directly use pre-trained CNNs
  - **Which** layer to use?

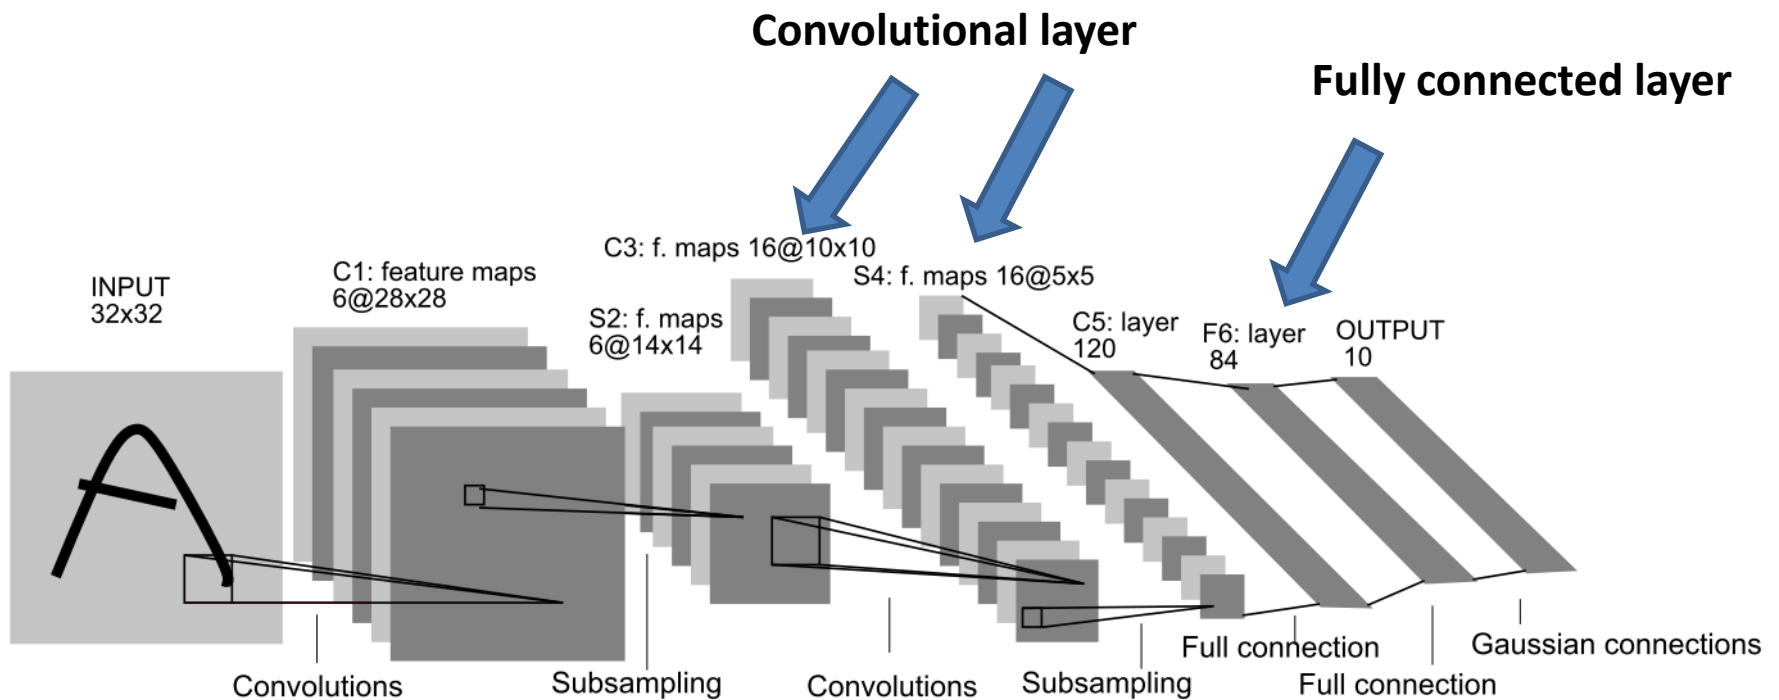


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.



# Deep Learning Model

- Directly use pre-trained CNNs
  - How to **pool** the features in a **convolutional** layer?

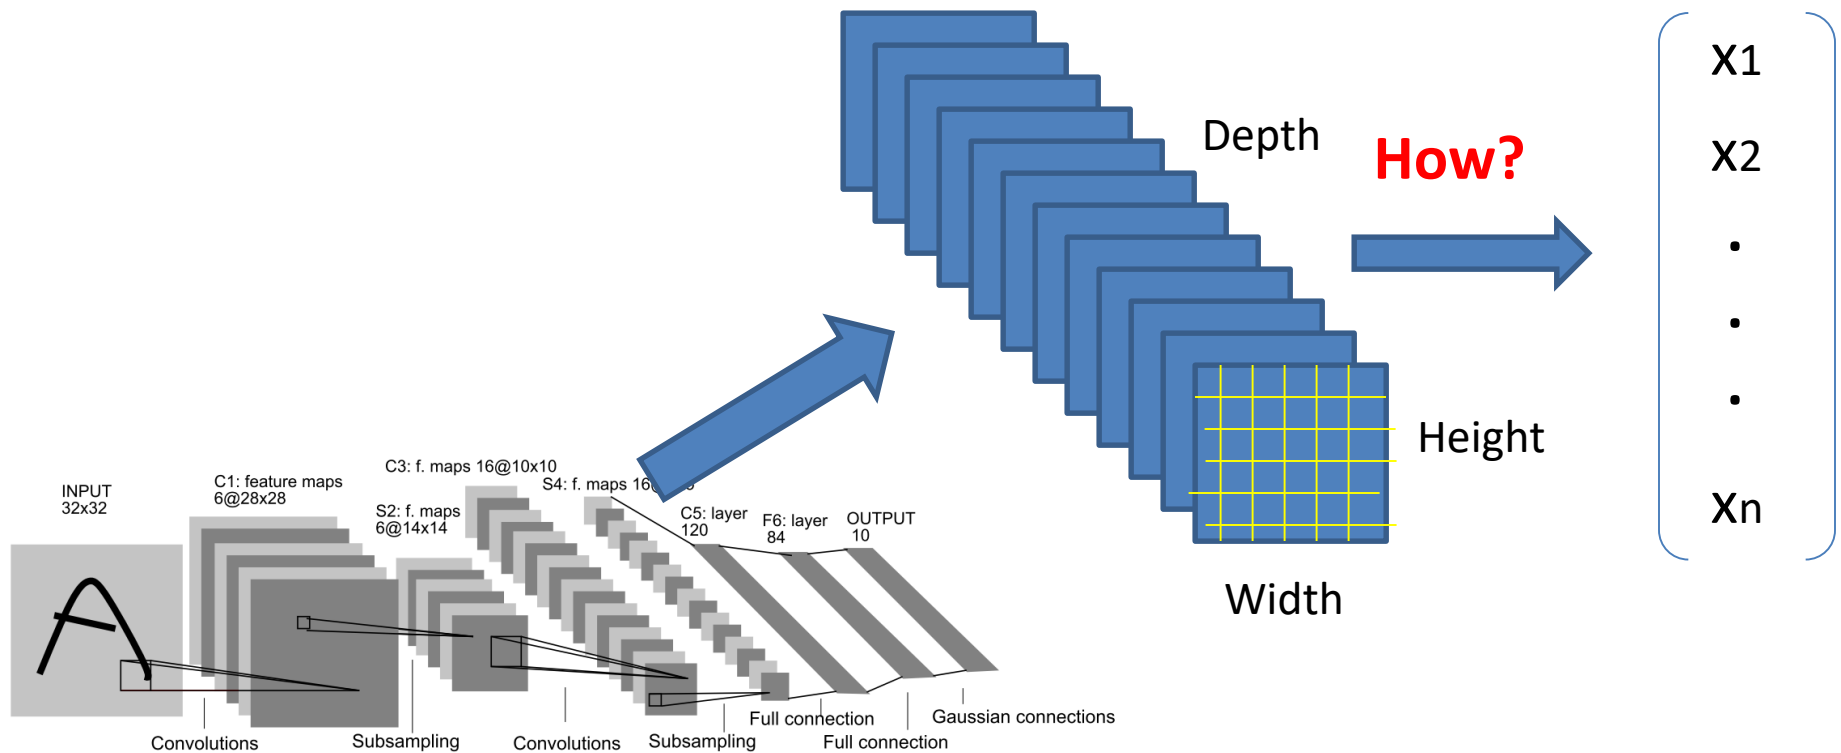



Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

# Deep Learning Model

- Directly use pre-trained CNNs
  - How to **pool** the features in a convolutional layer?



- Sum-pooling
- Max-pooling
- Grid-based max-pooling
- Region-based pooling
- Mixed sum & max pooling

# Deep Learning Model

- Fine-tune pre-trained CNNs
  - To incorporate extra information from the images of a new recognition task
  - Make the pre-trained CNNs adapt to this new task

## Pre-trained CNNs on



Image courtesy of Deng et al.

## Fine-tune





## New recognition task



<http://people.csail.mit.edu/bzhou/>

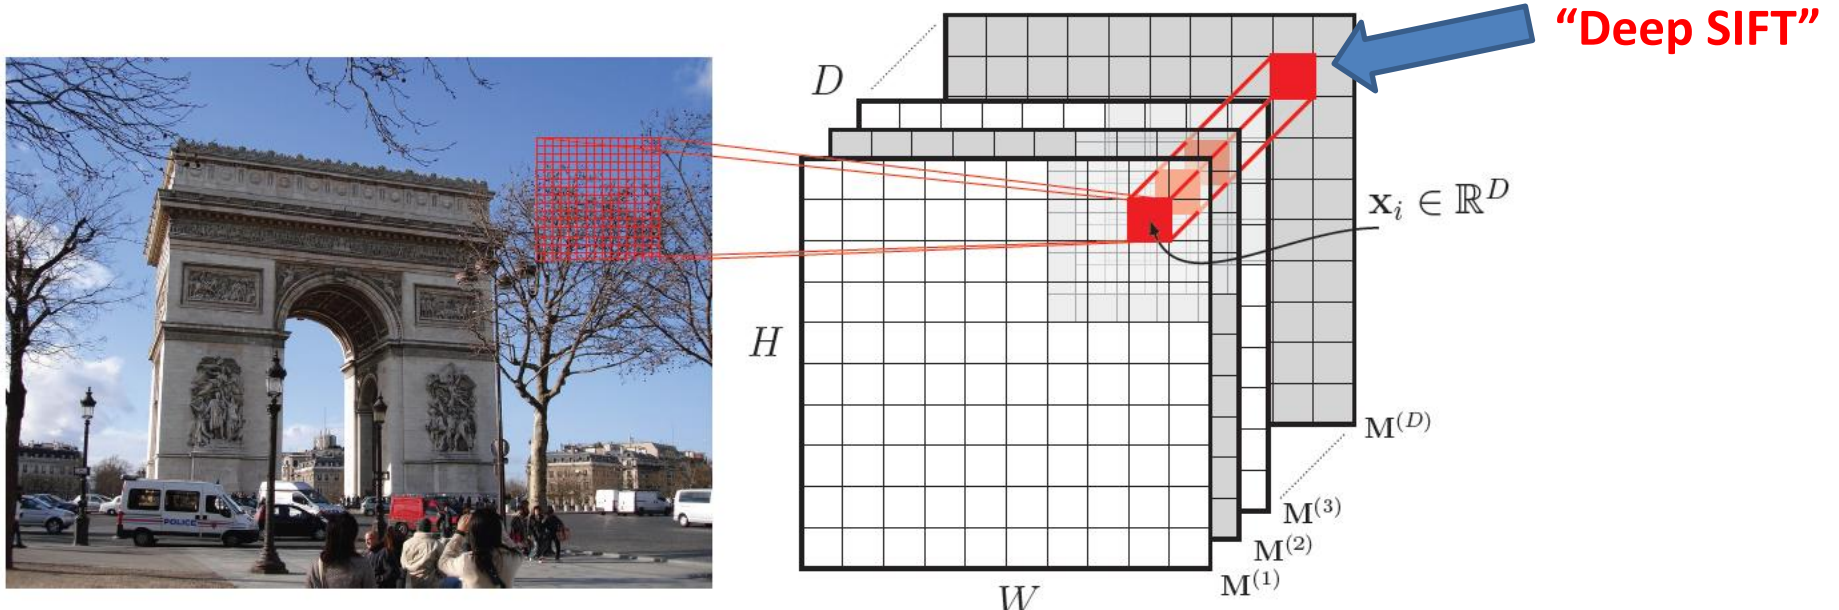
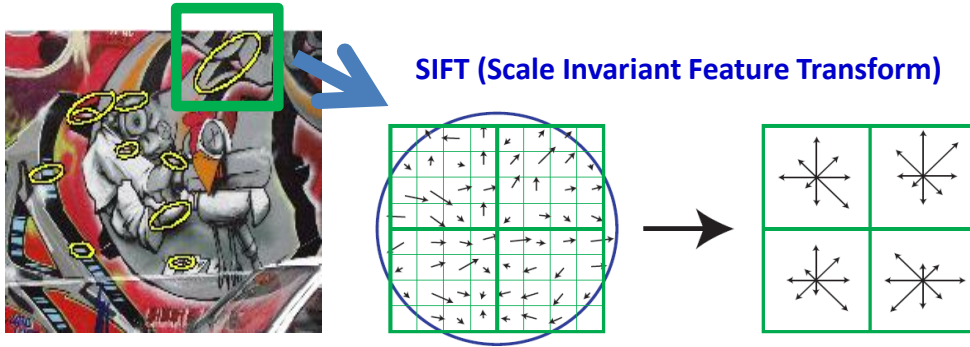
# Deep Learning Model

## BoVW model based on “Deep SIFT”



# Deep Learning Model




## BoVW model based on “Deep SIFT”



Source: Cao et. al, Where to Focus: Query Adaptive Matching for Instance Retrieval Using Convolutional Feature Maps




# Deep Learning Model



Or

**Visual word  
creation**



**Histogram  
generation**



**Image  
recognition**

# Deep Learning Package for R

- The support for performing deep learning with **R language** is increasing

Package	Multiple CPU	[Multiple] GPU	Cluster	Platforms
MXNetR	X	X		Linux\MacOS\ Windows
darch		X		Linux\MaxOS
H2o	X		X	Linux\MacOS\ Windows
deepnet				No information

Source: <http://www.rblog.uni-freiburg.de/2017/02/07/deep-learning-in-r/>

# Deep Learning Package for R

- MXNetR package ([link1](#), [link2](#), [Video](#))
  - R interface of the [MXNet](#) library (written in C++)
  - Contains feed-forward neural networks and [convolutional neural networks](#)
  - The [CPU](#) version can be easily installed in R

```
model <- mx.mlp(train.x, train.y, hidden_node=c(128,64),  
out_node=2, activation="relu",  
out_activation="softmax", num.round=100,  
array.batch.size=15, learning.rate=0.07, momentum=0.9,  
device=mx.cpu())  
  
preds = predict(model, testset)
```

# Deep Learning Package for R


- **MXNetR** package can visualise the networks

```
graph.viz(model$symbol$as.json())
```

- A helpful blog on deep learning with R

<https://www.r-bloggers.com/image-recognition-tutorial-in-r-using-deep-convolutional-neural-networks-mxnet-package/>

Computation graph



# Summary

- **Image analysis** is an important and broad area
- **Feature representation** is key for image analysis
- **Deep Learning** techniques are now widely used
- **Comparison** of deep learning software
- Issues to be resolved for image recognition
  - How to **transfer** the benefit of Deep Learning?
  - How to deal with **unsupervised** learning case?
  - How to better handle the **semantic gap**?
  - ...

