

CSCI446/946 Big Data Analytics

Week 12 Advanced Analytical Theory and
Methods: MapReduce and Hadoop

School of Computing and Information Technology
University of Wollongong Australia

Advanced Analytical Theory and Methods: MapReduce and Hadoop

- Overview
- Analytics for Unstructured Data
 - Use Cases, MapReduce, Hadoop
- The Hadoop Ecosystem
 - Pig, Hive, Hbase, Mahout
- NoSQL
- Summary



Overview

- The Apache Hadoop software library
 - A framework allows for the **distributed processing of large datasets** across clusters of computers using simple programming models
- Hadoop
 - How it stores data in a distributed system
 - How it implements a simple programming paradigm – **MapReduce**
 - How its tools help MapReduce programming

All the figures, tables and codes are from the book “[Data Science and Big Data Analytics: Discovering, Analyzing, Visualizing and Presenting Data](#)” unless indicated otherwise.

Analytics for Unstructured Data

- Unstructured data (text, image, video, etc.)
 - Data that has no inherent, consistent structure
- MapReduce paradigm offers the means to
 - break a large task into smaller tasks
 - run tasks in parallel
 - consolidate the outputs of the individual tasks into the final output
- Apache Hadoop implements MapReduce

Analytics for Unstructured Data

- Use Cases
 - [IBM Watson](#) won Jeopardy champions
 - [LinkedIn](#) utilizes Hadoop for multiple tasks
 - Process daily production database transaction logs
 - Examine the users' activities (views and clicks)
 - Feed the extracted data back to the production systems
 - [Yahoo!](#)'s Hadoop applications involve
 - Search index creation, web page content optimization, web ad placement optimization, spam filters, ...

Analytics for Unstructured Data

- MapReduce consists of two basic parts
 - A map step and a reduce step
- **Map** step
 - Applies an operation to a piece of data
 - Provides some **intermediate** output
- **Reduce** step
 - Consolidates the above **intermediate** outputs
 - Provides the final output

Analytics for Unstructured Data

- Each step uses key/value pairs as input and output, denoted as `<key, value>`
- The pairs can take `complex` forms
 - For example, the key is a filename, and the value is the entire content of the file
- A `simple illustration` of MapReduce
 - The task is to count the number of times each word appears in a collection of documents

Analytics for Unstructured Data

<1234, "For each word in each string">



Map

<For, 1> <each, 1> <word, 1> <in, 1> <each, 1> <string, 1>



Reduce

<For, 1>
<each, 2>
<word, 1>
<in, 1>
<string, 1>

- The map step parse the string into single words and emits a set of key/value pairs in the form <word, 1>;
- For each unique key, the reduce step sums the 1 values and outputs the <word, count> key/value pairs;

Analytics for Unstructured Data


- MapReduce
 - Has the **advantage** of being able to **distribute** the workload over a cluster of computers and run the tasks in **parallel**
 - The documents, or even their pieces, could be processed **simultaneously** during the map step
 - The processing of one portion of the input can be carried out **independently** of the others

Analytics for Unstructured Data

- MapReduce
 - A simple paradigm to understand, but **not easy to implement**, esp. in a distributed system
- Executing a MapReduce job requires
 - Management and coordination of activities of
 - Scheduling jobs, monitoring jobs
 - Spreading data, conducting the map step across
 - Collecting numerous intermediate outputs
 - Making the final output

Analytics for Unstructured Data

- Apache Hadoop
 - **Handles** these activities and more, and make most of them **transparent** to the users
 - An open source project managed and licenced by the Apache Software Foundation
- The development of MapReduce proceeds easily because
 - The MapReduce paradigm **has already been implemented** in Apache Hadoop



Analytics for Unstructured Data

- Apache Hadoop
 - Hadoop Distributed File System (HDFS)
 - How data is stored in a Hadoop environment
 - Structuring a MapReduce job in Hadoop
 - How a MapReduce job is run
 - Additional considerations
 - Developing a Hadoop MapReduce program

Analytics for Unstructured Data

- Hadoop Distributed File System (HDFS)
 - Provides the capability to **distribute** data across a cluster to use **parallel** processing of **MapReduce**
 - HDFS is **not an alternative** to common file systems, but depends on each disk drive's file system
 - HDFS breaks a file into blocks and stores the blocks across the clusters
 - The blocks of a file are stored on different machines
 - By default, creates three copies of each block across

Analytics for Unstructured Data

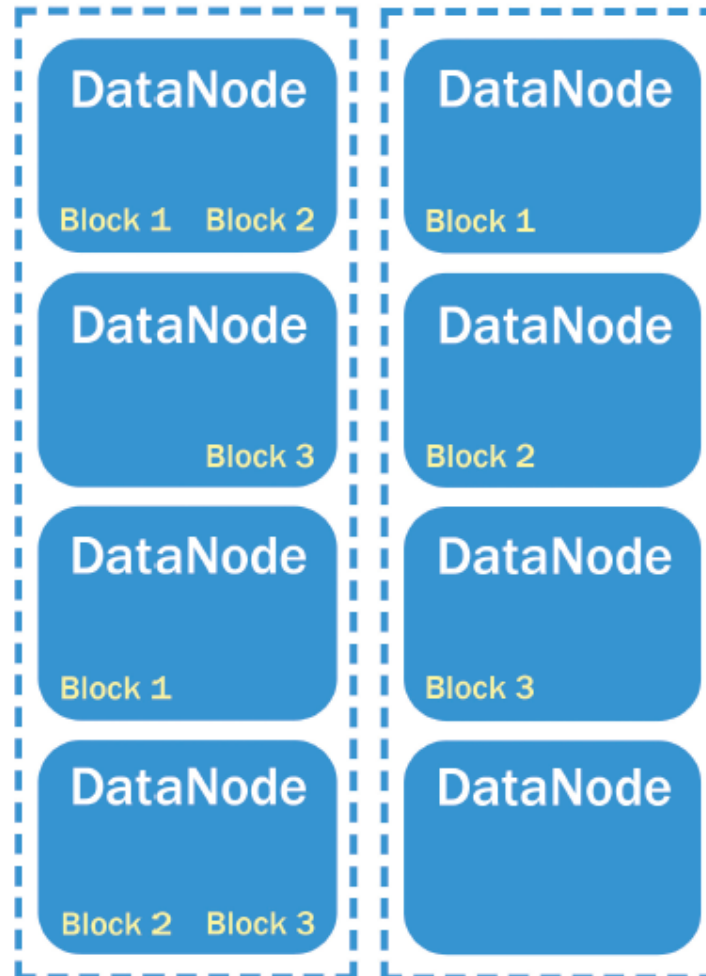
- HDFS uses **three** Java daemons (background processes)
 - **NameNode** determines and tracks where the blocks of a data file are stored. It runs on a single machine and resides in its memory
 - **DataNode** manages data stored on each machine
 - **Secondary NameNode** provides the capability to perform some of the NameNode tasks to reduce the load of NameNode
 - But it is not a backup or redundant NameNode

Analytics for Unstructured Data

Master Nodes



8 Worker Nodes across 2 Racks



Analytics for Unstructured Data

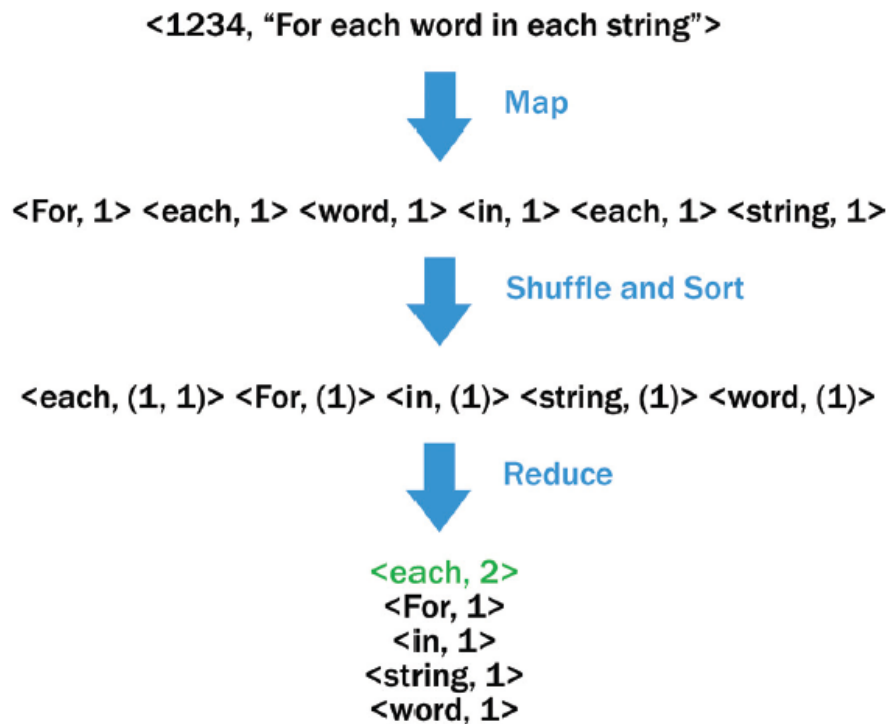
- A MapReduce program has **three classes**
 - The **driver class** provides details such as
 - Input and output file locations
 - Names of the mapper and reducer Java classes
 - Various job configuration options
 - The **mapper class** provides the logic to be processed on each data block corresponding to the specified input files
 - The **reducer class** provides the logic to process the values for each key to emit a key/value pair

Analytics for Unstructured Data

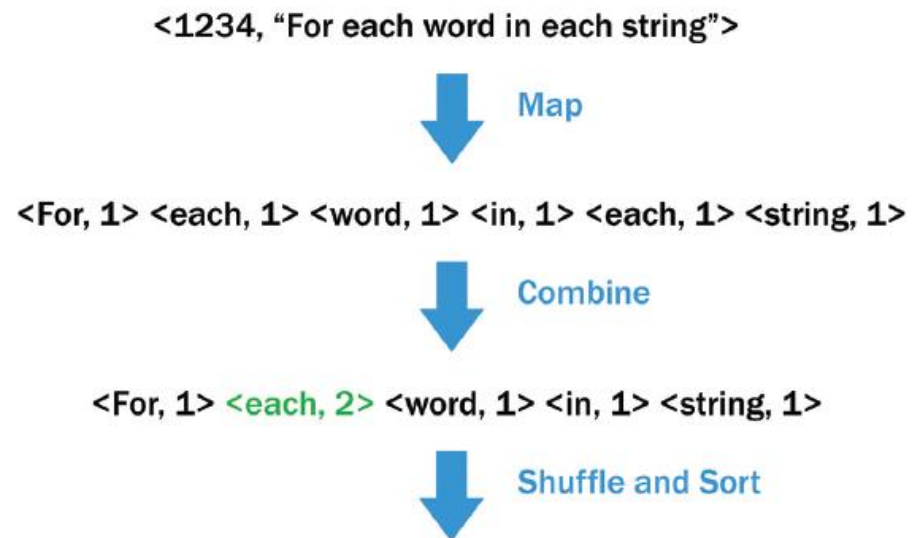
- Additional consideration in structuring a job
 - **Combiner** (“local reducer”)
 - It is a useful option to apply between the map task and the shuffle and sort
 - It applies the same logic used in the reducer on the output of each map task
 - **Partitioner**
 - It determines the reducers that receive keys and the corresponding list of values

Analytics for Unstructured Data

- Additional consideration in structuring a job



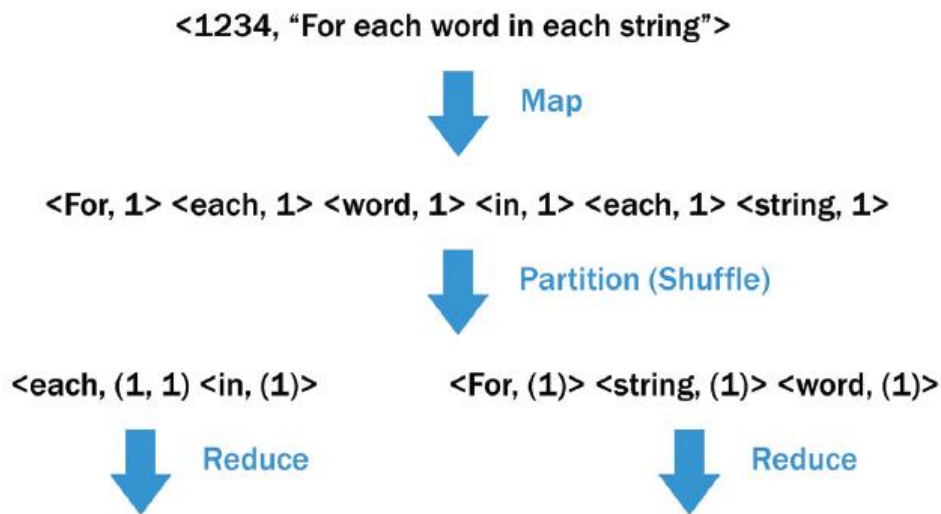
Shuffle and sort



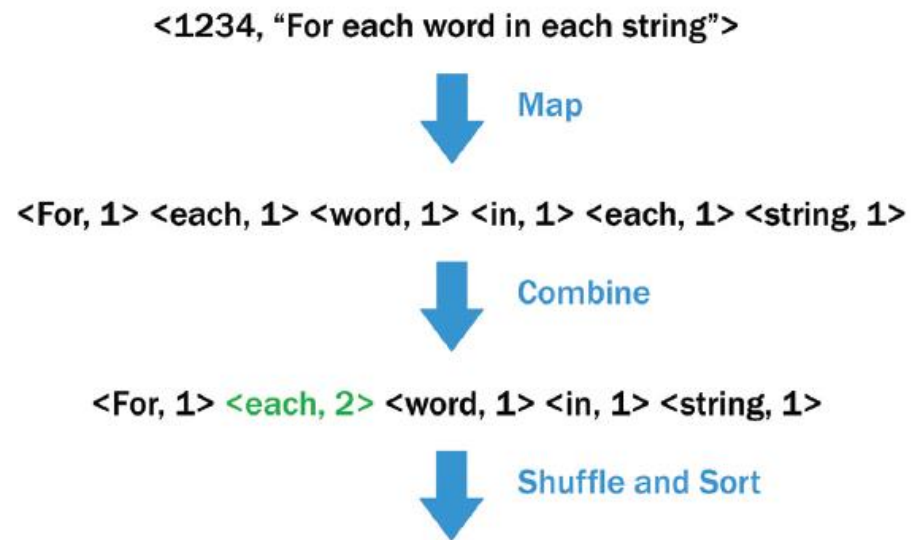
Combiner

Analytics for Unstructured Data

- Additional consideration in structuring a job



Partitioner



Combiner

Analytics for Unstructured Data

- Developing and executing a program (via Java)
 - A typical MapReduce program consists of **three Java files**, one each for the driver code, map code, and reduce code
 - Java files can be written for the **combiner** and **partitioner**
 - Java code is compiled as a Java Archive (JAR) file
 - The **JAR file** is executed against the specified **HDFS input files**

Analytics for Unstructured Data

- Three key challenges to a new developer
 - Defining the logic of the code to use the MapReduce paradigm
 - Learning the Apache Hadoop Java classes, methods, and interfaces
 - Implementing the driver, map, and reduce functionality in Java
- Some prior experience with Java will help

Analytics for Unstructured Data

- Three **key challenges** to a new developer
 - **Defining** the logic of the code to use the MapReduce paradigm
 - **Learning** the Apache Hadoop Java classes, methods, and interfaces
 - **Implementing** the driver, map, and reduce functionality in Java
- Some prior experience with Java will help

Analytics for Unstructured Data

- For non-Java developers
 - Hadoop Streaming API
 - Hadoop-streaming.jar file: accepts the HDFS paths for the input/output files and the paths for the files that implement the map and reduce functionality
 - There are some important considerations/limitations
 - Hadoop pipes
 - A mechanism uses compiled C++ code for the map and reduce functionality
 - An advantage of using C++ is the numerical libraries

Analytics for Unstructured Data

- To directly work with data in HDFS
 - Use the C API (libhdfs) or the Java API provided with Apache Hadoop
 - These APIs allow reads and writes to HDFS data files outside the typical MapReduce paradigm
 - Such approach may be helpful when
 - Debug a MapReduce job by examining the input data
 - Transform the HDFS prior to running a MapReduce job

Analytics for Unstructured Data

- MapReduce 2.0
 - Also called Yet Another Resource Negotiator (YARN)
- The two functionalities are separated
 - MapReduce functionality
 - Those to manage job running and associated responsibilities in a distributed environment
 - YARN makes it possible for utilizing the paradigms other than MapReduce in Hadoop

The Hadoop Ecosystem

- Hadoop's proprietary and open source tools
 - Make Apache Hadoop easier to use
 - Provide additional functionality and features
- Hadoop-related Apache projects
 - **Pig**: A high-level data-flow programming language
 - **Hive**: Provide SQL-like access
 - **Mahout**: Provide analytical tools
 - **Hbase**: Provides real-time reads and writes

The Hadoop Ecosystem

- Apache Pig
 - Consists of a data flow language, Pig Latin, and an environment to execute the Pig code
 - Benefit is simplifying the tasks of developing and executing a MapReduce job
 - When Pig commands are executed, the running of a job at background is transparent to the user
 - Three main characteristics
 - Ease of programming, behind-the-scenes code optimization, and extensibility of capabilities

The Hadoop Ecosystem

- Apache Hive
 - Similar to Pig, Hive enables users to process data **without explicitly writing** MapReduce code
 - One key difference to Pig
 - Hive language (HiveQL) resembles **Structured Query Language (SQL)** rather than a script language
 - Hive may be a **good tool** to use if
 - Data easily fits into a table structure
 - Data is already in HDFS
 - Developers are comfortable with SQL

The Hadoop Ecosystem

- Apache HBase

- Pig and Hive are intended for batch applications
- Differently, Hbase provides real-time read and write access to large-scale datasets
- Is built upon HDFS
- Share the workload over a large number of nodes in a distributed cluster
- Uses a key/value structure to store the contents of an HBase table

The Hadoop Ecosystem

- Apache Mahout
 - Tools such as R may suffer from performance issues with the large datasets in stored in Hadoop
 - Supports the application of analytical techniques within the Hadoop environment
 - Provides Java libraries to apply analytical techniques in a scalable manner to Big Data
 - Implemented algorithms
 - Classification, clustering, collaborative filtering

NoSQL

- **NoSQL** (Not only Structured Language)
 - A term used to describe those **data stores** that are applied to **unstructured** data
 - As the size of data grows, the solution can scale by adding machines to the distributed system
 - Four major categories of NoSQL tools
 - Key/value store, Document store, Column family store, Graph databases
 - The choice of data store is take-dependent

Summary

- The **MapReduce** paradigm and its implementation in Apache **Hadoop**
- The Hadoop Distributed File System (**HDFS**)
- Developing a Hadoop MapReduce program
- Apache projects in the Hadoop Ecosystem



